A survey of generic IND-CCA2 transformations

Ganyu Xu

July 11, 2024

1 Preliminaries

1.1 Public-key encryption schemes

A public key encryption scheme $\mathtt{PKE} = (\mathtt{KeyGen}, \mathtt{E}, \mathtt{D})$ is a collection of three routines. $\mathtt{KeyGen}(1^\lambda)$ takes the security parameter as input and returns a keypair $(\mathtt{pk}, \mathtt{sk})$. $\mathtt{E}(\mathtt{pk}, m)$ takes some public key and some plaintext message $m \in \mathcal{M}_{\mathtt{PKE}}$ and output a ciphertext $c \in \mathcal{C}_{\mathtt{PKE}}$. Where the encryption routine is probabilistic, we model the randomness using a coin $r \in \mathcal{R}$ such that $E(\mathtt{pk}, m; r)$ is deterministic with an explicit r. Finally, $\mathtt{D}(\mathtt{sk}, c)$ uses the secret key to decrypt the ciphertext.

1.1.1 Correctness

Conventionally we require a $\mathtt{PKE}$ to be perfectly correct. This means that for all possible key pairs $(\mathtt{pk}, \mathtt{sk})$ and plaintexts m, the decryption routine always correctly inverts the encryption routine: $\mathtt{D}(\mathtt{sk}, \mathtt{E}(\mathtt{pk}, m)) = m$.

Where perfect correctness is not achieved, such as with most lattice-based encryption schemes, we need to account for the possiblity that decryption can fail. If under some keypair $(\mathtt{pk}, \mathtt{sk})$, a plaintext-ciphertext pair (m, c) is such that $c \xleftarrow{\$} \mathtt{E}(\mathtt{pk}, m)$ is obtained from encrypting m (probabilistically) but $m \neq \mathtt{D}(\mathtt{sk}, c)$, we call it a decryption failure. For probabilistic encryption routines where the coin is uniformly sampled from the coin space, we can quantify the probability that m triggers a decryption failure. From here, we can take the distribution of all keypairs and quantify the "correctness" of a (possibly imperfectly correct) encryption scheme. The following definition of δ-correctness is directly taken from [BDK$^+$18].

Definition 1.1 (δ-correctness). *A probabilistic public-key encryption scheme $PKE = (KeyGen, E, D)$ is δ-correct if the expected maximal probability of decryption failure taken across the distribution of keypairs is at most δ:*

$$E \left[\max_{m \in \mathcal{M}} P \left[D(sk, E(pk, m)) \right] \right] \leq \delta$$

where the expectation is taken over the distribution of keypairs and the probability is taken over the distribution of coins.

[HHK17] also defined an adversarial game in which the adversary's goal is to find some plaintext message to trigger a decryption failure. This adversarial game meaningfully models the real-world scenario in which decryption failure can reveal information about the secret key. Notice that when evaluating the win condition, the coin is uniformly random instead of being chosen by the adversary.

Algorithm 1 `CORS`

1: $(\mathtt{pk}, \mathtt{sk}) \xleftarrow{\$} \mathtt{KeyGen}(1^\lambda)$

2: $m \xleftarrow{\$} A_{\mathtt{CORS}}(1^\lambda, \mathtt{pk}, \mathtt{sk})$

3: **return** $\llbracket \mathtt{D}(\mathtt{sk}, \mathtt{E}(\mathtt{pk}, m)) \neq m \rrbracket$

Figure 1: The correctness game `CORS`

The definition of δ-correctness sets an explicit upper bound on the probability that any plaintext triggers decryption failure. This means that even if the `CORS` adversary actually finds the message m that is the most likely to trigger decryption failure, the probability of winning the `CORS` game is still upper-bounded by δ.

Lemma 1.0.1. *If `PKE` is δ-correct, then for all `CORS` adversaries A, even computationally unbounded ones, the probability of winning the `CORS` game is at most δ*

The values of δ for Kyber are taken directly from [ABD+19]

security level	δ
Kyber512	2^{-139}
Kyber768	2^{-164}
Kyber1024	2^{-174}

Table 1: Concrete δ for Kyber

1.1.2 Security

An one-way adversary $A = (A_1, A_2)$ consists of two sub-routines A_1 and A_2. $s \xleftarrow{\$} A_1^{\mathcal{O}_1}(1^\lambda, \text{pk})$ takes the security parameter, some public-key, access to some oracle(s) $\mathcal{O}_1$, and outputs some intermediate state s. $\hat{m} \leftarrow A_2^{\mathcal{O}_2}(1^\lambda, \text{pk}, c^*)$ resumes from the output of A_1 and takes some challenge ciphertext, then tries to guess the corresponding decryption.

The advantage $\text{Adv}_{\text{OW-ATK}}(A)$ of an `OW-ATK` adversary is the probability that its guess is correct.

Definition 1.2. *A `PKE` is `OW-ATK` secure if for all efficient adversaries A, the advantage in the `OW-ATK` game is neligigble with respect to the security parameter:*

$$\textit{Adv}_{\textit{OW-ATK}}(A) \leq \text{negl}(\lambda)$$

An **indistinguishability** adversary $A = (A_1, A_2)$ similarly consists of two sub-routines. The first sub-routine adversarially chooses two distinct plaintext messages, and the second sub-routine tries to distinguish which of the two plaintext messages is the decryptino of the challenge encryption. The advantage of an indistinguishability adversary is defined by $\text{Adv}_{\text{IND-ATK}}(A) = P[\hat{b} = b] - \frac{1}{2}$

Definition 1.3. *A `PKE` is `IND-ATK` secure if for all efficient adversaries A, the advantage in the indistinguishability game is negligible with respect to the security parameter*

$$\textit{Adv}_{\textit{IND-ATK}}(A) \leq \text{negl}(\lambda)$$

The security games are described in details in figure 2

Algorithm 2 OW-ATK game
1: $(\text{pk}, \text{sk}) \xleftarrow{\$} \text{KeyGen}(1^\lambda)$
2: $s \xleftarrow{\$} A_1^{\mathcal{O}_1}(1^\lambda, \text{pk})$
3: $m^* \xleftarrow{\$} \mathcal{M}$
4: $c^* \xleftarrow{\$} \text{E}(\text{pk}, m^*)$
5: $\hat{m} \leftarrow A_2^{\mathcal{O}_2}(1^\lambda, \text{pk}, s, c^*)$
6: **return** $[\![\hat{m} = m^*]\!]$

Algorithm 3 IND-ATK game
1: $(\text{pk}, \text{sk}) \xleftarrow{\$} \text{KeyGen}(1^\lambda)$
2: $(m_0, m_1) \xleftarrow{\$} A_1^{\mathcal{O}_1}(1^\lambda, \text{pk})$
3: $b \xleftarrow{\$} \{0, 1\}$
4: $c^* \xleftarrow{\$} \text{E}(\text{pk}, m_b)$
5: $\hat{b} \leftarrow A_2^{\mathcal{O}_2}(1^\lambda, \text{pk}, s, c^*)$
6: **return** $[\![\hat{b} = b]\!]$

Figure 2: The indistinguishability security game

The capabilities of the adveraries are modeled using different collections of oracles. In standard security requirements, adversaries with access to no additional oracles can only mount chosen plaintex attacks (CPA),

adversaries with access to decryption oracle $\mathcal{O}^{\mathrm{D}}$ only before the receiving challenge ciphertext can mount non-adaptive chosen ciphertext attacks (CCA1), adversaries with access to decryption oracle both before and after receiving the challenge ciphertext can mount adaptive chosen ciphertext attacks (CCA2).

[HHK17] also defined two non-standard oracles and the corresponding attacks. The plaintext checking oracle $\mathtt{PCO}(m, c)$ returns 1 if m is a decryption of c and 0 otherwise. The ciphertext validation oracle $\mathtt{CVO}(c)$ returns 1 if c is a valid ciphertext and 0 otherwise.

Algorithm 4 $\mathtt{PCO}(m, c)$	**Algorithm 5** $\mathtt{CVO}(c)$
return $[\![\mathrm{D}(\mathtt{sk}, c) = m]\!]$	1: **return** $[\![\mathrm{D}(\mathtt{sk}, c) \in \mathcal{M}]\!]$

Figure 3: $\mathtt{PCO}$ and $\mathtt{CVO}$

Here is an overview of the various kinds of attacks and their associated oracles

ATK	$\mathcal{O}_1$	$\mathcal{O}_2$
CPA	—	
CCA1	$\mathcal{O}^{\mathrm{D}}$	-
CCA2	$\mathcal{O}^{\mathrm{D}}$	
PCVA	PCO, CVO	
PCA	PCO	
VA	CVO	

Table 2: Attacks and associated oracle access

[HHK17] stated a "well-known" result that the $\mathtt{IND\text{-}CPA}$ security of a scheme with a large message space implies $\mathtt{OW\text{-}CPA}$ security:

Theorem 1.1. *For every $\mathtt{IND\text{-}CPA}$ adversary B against some $\mathtt{PKE}$, there exists an $\mathtt{OW\text{-}CPA}$ adversary A against the same $\mathtt{PKE}$ such that:*

$$\boldsymbol{Adv}_{\mathtt{OW\text{-}CPA}}(A) = \frac{1}{|\mathcal{M}|} + \boldsymbol{Adv}_{\mathtt{IND\text{-}CPA}}(B)$$

1.1.3 Spread and rigidity

The spread of a public key encryption scheme measures the diffusion the encryption routine's output. The higher the spread, the lower the probability of obtaining any specific ciphertext.

Definition 1.4 (γ-spread)**.** *For a given keypair $(\boldsymbol{pk}, \boldsymbol{sk})$ and plaintext message $m \in \mathcal{M}$, the* min-entropy *of the encryption routine is:*

$$\text{min-entropy}(\boldsymbol{pk}, m) := -\log_2 \left(\max_{c \in \mathcal{C}} P\left[c = \boldsymbol{E}(\boldsymbol{pk}, m) \right] \right)$$

A $\mathtt{PKE}$ has γ-spread if for all keypairs $(\boldsymbol{pk}, \boldsymbol{sk})$ and plaintext $m \in \mathcal{M}$:

$$\text{min-entropy}(\boldsymbol{pk}, m) \leq \gamma$$

Having γ sperad means that for any keypair $(\mathtt{pk}, \mathtt{sk})$, plaintext m, and ciphertext c:

$$P\left[c = \mathtt{E}(\mathtt{pk}, m) \right] \leq 2^{-\lambda}$$

Finally, *rigidity* conveys the idea that a ciphertext cannot be perturbed without becoming either invalid or decrypting to another plaintext

Definition 1.5 (rigidity)**.** *$\mathtt{PKE}(\boldsymbol{KeyGen}, \boldsymbol{E}, \boldsymbol{D})$ is* rigid *if for all keypairs $(\boldsymbol{pk}, \boldsymbol{sk})$ and ciphertext c, either $\boldsymbol{D}(\boldsymbol{sk}, c) = \bot$ or $\boldsymbol{E}(\boldsymbol{pk}, \boldsymbol{D}(\boldsymbol{sk}, c)) = c$*

1.2 Key encapsulation mechanism

A key encapsulation mechanism $\texttt{KEM} = (\texttt{KeyGen}, \texttt{Encap}, \texttt{Decap})$ is a collection of three routines. The key generation routine $(\texttt{pk}, \texttt{sk}) \xleftarrow{\$} \texttt{KeyGen}(1^\lambda)$ takes the security parameter 1^λ and returns a keypair. The encapsulation routine $(c, K) \xleftarrow{\$} \texttt{Encap}(\texttt{pk})$ takes the public key and outputs some ciphertext $c \in \mathcal{C}_{\texttt{KEM}}$ and some shared secret $K \in \mathcal{K}_{\texttt{KEM}}$. Finally, the decapsulation routine $K \leftarrow \texttt{Decap}(\texttt{sk}, c)$ takes the secret key and a ciphertext and outputs the correponding shared secret.

Correctness: similar to a $\texttt{PKE}$, key encapsulation mechanisms are usually required to be perfectly correct, meaning that for all keypairs $(\texttt{pk}, \texttt{sk})$, decapsulation always outputs the same shared secret as the encapsulation

$$\mathrm{P}\left[(c, K_1) \xleftarrow{\$} \texttt{Encap}(\texttt{pk}); K_2 \leftarrow \texttt{Decap}(\texttt{sk}, c); K_1 = K_2\right] = 1$$

However, where decapsulation failure has a non-zero probability, we need to upperbound this quantity.

Definition 1.6 (δ-correctness)**.** *A* $\texttt{KEM} = (\texttt{KeyGen}, \texttt{Encap}, \texttt{Decap})$ *is δ-correct if the probability of decapsulation failure taken across the keypair distribution is at most δ:*

$$P\left[\texttt{Decap}(\texttt{sk}, c) \neq K \mid (\texttt{pk}, \texttt{sk}) \xleftarrow{\$} \texttt{KeyGen}(); (c, K) \xleftarrow{\$} \texttt{Encap}(\texttt{pk})\right] \leq \delta$$

Security: the security of a $\texttt{KEM}$ is defined in an adversarial game in which the adversary's goal is to distinguish between shared secret derived from encapsulation and uniformly random samples. An adversary $A = (A_1, A_2)$ contains two sub-routines with access to some oracle $\mathcal{O}$ depending on game.

Algorithm 6 IND-CCA2 game

1: $(\texttt{pk}, \texttt{sk}) \xleftarrow{\$} \texttt{KeyGen}(1^\lambda)$
2: $s \xleftarrow{\$} A_1^{\mathcal{O}^{\texttt{Decap}}}(1^\lambda, \texttt{pk})$
3: $(c^*, K_0) \xleftarrow{\$} \texttt{Encap}(\texttt{pk})$
4: $K_1 \xleftarrow{\$} \mathcal{K}_{\texttt{KEM}}$
5: $b \xleftarrow{\$} \{0, 1\}$
6: $\hat{b} \xleftarrow{\$} A_2^{\mathcal{O}^{\texttt{Decap}}}(1^\lambda, \texttt{pk}, s, c^*, K_b)$
7: **return** $[\![\hat{b} = b]\!]$

Algorithm 7 Decap oracle $\mathcal{O}^{\texttt{Decap}}(c \neq c^*)$

1: **return** $\texttt{Decap}(\texttt{sk}, c)$

Figure 4: The IND-CCA2 game for KEM

Definition 1.7 (IND-CCA2 security)**.** *A* $\texttt{KEM} = (\texttt{KeyGen}, \texttt{Encap}, \texttt{Decap})$ *is IND-CCA2 secure if no efficient adversary has non-negligible advantage in the IND-CCA2 game.*

2 Modular Fujisaki-Okamoto transformation

The Fujisaki-Okamoto transformation [FO99] and its KEM variations [HHK17] achieve security against adaptive chosen-ciphertext attacks through *de-randomization* and *re-encryption*. In particular, the modular FO transformation contains two steps. The first step (denoted the T transformation) takes a $\texttt{PKE} = (\texttt{KeyGen}, \texttt{E}, \texttt{D})$ and outputs a $\texttt{PKE}_1 = (\texttt{KeyGen}, \texttt{E}_1, \texttt{D}_1)$. The key generation remains unchanged, but the encryption routine is *de-randomized* by deriving the coin from the plaintext using a hash function $G : \mathcal{M}_{\texttt{PKE}} \to \mathcal{R}_{\texttt{PKE}}$, and the decryption routine uses *re-encryption* to reject invalid ciphertexts.

Algorithm 8 $\mathrm{E}_1(\mathrm{pk}, m)$	**Algorithm 9** $\mathrm{D}_1(\mathrm{sk}, c)$
1: $r \leftarrow G(m)$	1: $\hat{m} \leftarrow D(\mathrm{sk}, c)$
2: $c \leftarrow \mathrm{E}(\mathrm{pk}, m; r)$	2: **if** $\hat{m} \in \mathcal{M}_{\mathrm{PKE}} \wedge \mathrm{E}(\mathrm{pk}, \hat{m}; G(\hat{m})) = c$ **then**
3: **return** c	3: **return** $\hat{m}$
	4: **end if**
	5: **return** $\perp$

Figure 5: T transformation

[HHK17] states the security property of PKE$_1$ as follows:

Theorem 2.1. *If PKE is δ-correct and has γ spread, then for every* OW-PCVA *adversary B against PKE$_1$ who makes q_G hash queries, q_V ciphertext validation queries, and q_P plaintext checking queries, there exists an* OW-CPA *adversary A such that:*

$$\textit{Adv}(B) \leq q_V \cdot 2^{-\gamma} + (q_G + q_P) \cdot \delta + (1 + q_G + q_P) \cdot \textit{Adv}(A)$$

In other words, if PKE is OW-CPA secure, then PKE$_1$ is OW-PCVA secure, though the security is non-tight. On the other hand, if PKE is IND-CPA secure, the OW-PCVA security is tight:

Theorem 2.2. *for every* OW-PCVA *adversary B against PKE$_1$ who makes q_G hash queries, q_V ciphertext validation queries, and q_P plaintext checking queries, there exists an* IND-CPA *adversary A such that:*

$$\textit{Adv}(B) \leq q_V \cdot 2^{-\gamma} + (q_G + q_P) \cdot \delta + \frac{2 \cdot q_G + 1}{|\mathcal{M}_{PKE}|} + 3 \cdot \textit{Adv}(A)$$

Proof. We will provide a sketch of proof for theorem 2.1 and 2.2. The main idea is to construct an OW-CPA/IND-CPA adversary A who can simulate the OW-PCVA game and use the OW-PCVA adversary B as a sub-routine. However, there are three main difficulties with constructing a "convincing" simulation:

1. A has no access to PCO

2. A has no access to CVO

3. The challenge ciphertext A receives is obtained using a truly random coin, but the challenge ciphertext B expects is obtained using a pseudorandom coin $r \leftarrow G(m)$

$\square$

2.1 From PKE to KEM

The second part of the modular FO transform (denoted the U transform) takes a PKE and outputs a KEM. Depending on whether the input PKE is rigid and whether the output KEM rejects invalid ciphertext implicitly or explicitly, the security requirements for the input PKE and the construction of the KEM will vary. [HHK17] presents four variations of the U transform, which are listed in table 3

Name	PKE requirements	rejection	shared secret
$U^{\perp}$	OW-PCVA	$\perp$	$H(m, c)$
$U^{\not\perp}$	OW-PCA	$H(s, c)$	$H(m, c)$
$U_m^{\perp}$	rigidity + OW-VA	$\perp$	$H(m)$
$U_m^{\not\perp}$	rigidity + OW-CPA	$H(s, c)$	$H(m)$

Table 3: A summary of variants of U transformations

For the remaining of this section we will present the four transformations and sketch proofs of their securities.

2.1.1 KEM with explicit rejection from randomized PKE

Let $\mathtt{PKE} = (\mathtt{KeyGen}, \mathtt{E}, \mathtt{D})$ be a public-key encryption scheme, and $H : \mathcal{M}_{\mathtt{PKE}} \times \mathcal{C}_{\mathtt{PKE}} \to \mathcal{K}_{\mathtt{KEM}}$ be a hash function. The $U^{\perp}$ transformation outputs a $\mathtt{KEM}^{\perp} = (\mathtt{KeyGen}^{\perp}, \mathtt{Encap}^{\perp}, \mathtt{Decap}^{\perp})$, where the key generation routine remains unchanged: $\mathtt{KeyGen}^{\perp} = \mathtt{KeyGen}$. The encapsulation and decapsulation routines are described in figure 6

Algorithm 10 $\mathtt{Encap}^{\perp}(\mathtt{pk})$

1: $m \xleftarrow{\$} \mathcal{M}_{\mathtt{PKE}}$
2: $c \xleftarrow{\$} \mathtt{E}(\mathtt{pk}, m)$
3: $K \leftarrow H(m, c)$
4: **return** (c, K)

Algorithm 11 $\mathtt{Decap}^{\perp}(\mathtt{sk}, c)$

1: $\hat{m} \leftarrow \mathtt{D}(\mathtt{sk}, c)$
2: **if** $\hat{m} \in \mathcal{M}_{\mathtt{PKE}}$ **then**
3: $\quad K \leftarrow H(\hat{m}, c)$
4: $\quad$ **return** K
5: **end if**
6: **return** $\perp$

Figure 6: $U^{\perp}$ routines

Under the random oracle model, the security of $\mathtt{KEM}^{\perp}$ depends on the security of the input $\mathtt{PKE}$.

Theorem 2.3. *For every IND-CCA2 adversary B against $\mathtt{KEM}^{\perp}$, there exists an OW-PCVA adversary A against the underlying PKE such that*

$$\mathit{Adv}(B) \leq \mathit{Adv}(A)$$

Proof. For a sketch of proof, we argue that under the random oracle model, $H(m, c)$ is indistinguishable from a uniformly random sample from $\mathcal{K}_{\mathtt{KEM}}$ from the adversary B's perspective unless B somehow queries H on (m, c). This means that an OW-PCVA adversary can perfectly simulate the decapsulation oracle using uniformly random samples, as long as the outputs from the simulated decapsulation oracle are consistent with the outputs from the hash functions.

The only flaw in the simulated decapsulated oracle lies in the case when B queries $H(m^*, c^*)$, but because A has access to PCO, A can detect that B has successfully recovered m^* and uses the recovered m^* to win the OW-PCVA game. In other words, A keeps running B in a simulation until B makes the special hash query, at which time A simply terminates B and outputs the answer.

To make the simulated decapsulation oracle and the hash oracle consistent, A maintains two tapes $\mathcal{L}^{\mathtt{Decap}}$ and $\mathcal{L}^{H}$ that record the queries made to the decapsulation oracle and the hash oracles respectively:

$$(c, K) \in \mathcal{L}^{\mathtt{Decap}} \Leftrightarrow K = \mathcal{O}^{\mathtt{Decap}}(c)$$
$$(m, c, K) \in \mathcal{L}^{H} \Leftrightarrow K = H(m, c)$$

If $\mathcal{O}^{\mathtt{Decap}}$ has been queried with some valid ciphertext c, then when H is queried with a corresponding input (m, c) where m is the decryption of c, H needs to output the same value as $\mathcal{O}^{\mathtt{Decap}}$. Similarly, if H has been queried with some valid plaintext-ciphertext pair (m, c), then when $\mathcal{O}^{\mathtt{Decap}}$ is queried with c, the two oracles need to output the same value. The details of the simulation are described in figure 7

Algorithm 12 $\mathcal{O}_1^{\texttt{Decap}}(c)$

1: **if** $\exists (\tilde{c}, \tilde{K}) \in \mathcal{L}^{\texttt{Decap}} : \tilde{c} = c$ **then**
2:      **return** $\tilde{K}$
3: **end if**
4: **if** $\texttt{CVO}(c) = 1$ **then**
5:      $K \xleftarrow{\$} \mathcal{K}_{\texttt{KEM}}$
6:      Append (c, K) to $\mathcal{L}^{\texttt{Decap}}$
7:      **return** K
8: **end if**
9: **return** $\bot$

Algorithm 13 $H_1(m, c)$

1: **if** $\exists (\tilde{m}, \tilde{c}, \tilde{K}) \in \mathcal{L}^H : (\tilde{m}, \tilde{c}) = (m, c)$ **then**
2:      **return** $\tilde{K}$
3: **end if**
4: $K \xleftarrow{\$} \mathcal{K}_{\texttt{KEM}}$
5: **if** $\texttt{PCO}(m, c) = 1$ **then**
6:      **if** $\exists (\tilde{c}, \tilde{K}) \in \mathcal{L}^{\texttt{Decap}} : \tilde{(c)} = c$ **then**
7:          **return** $\tilde{K}$
8:      **else**
9:          Append (c, K) to $\mathcal{L}^{\texttt{Decap}}$
10:      **end if**
11: **end if**
12: Append (m, c, K) to $\mathcal{L}^H$
13: **return** K

Figure 7: Patched oracles in $U^\perp$

$\square$

2.1.2   KEM with implicit rejection from randomized PKE

Given $\texttt{PKE} = (\texttt{KeyGen}, \texttt{E}, \texttt{D})$, the $U^{\not\perp}$ transformation outputs $\texttt{KEM} = (\texttt{KeyGen}^{\not\perp}, \texttt{Encap}^{\not\perp}, \texttt{Decap}^{\not\perp})$. With implicit rejection, decapsulation routine returns some output from H even when the input ciphertext c is malformed. Specifically, the implicit rejection value depends on the ciphertext and some secret value s that is a part of the secret key.

Algorithm 14 $\texttt{KeyGen}^{\not\perp}$

1: $(\texttt{pk}, \texttt{sk}') \xleftarrow{\$} \texttt{KeyGen}(1^\lambda)$
2: $s \xleftarrow{\$} \mathcal{M}_{\texttt{PKE}}$
3: $\texttt{sk} = (\texttt{sk}', s)$
4: **return** $(\texttt{pk}, \texttt{sk})$

Algorithm 15 $\texttt{Encap}^{\not\perp}(\texttt{pk})$

1: $m \xleftarrow{\$} \mathcal{M}_{\texttt{PKE}}$
2: $c \xleftarrow{\$} \texttt{E}(\texttt{pk}, m)$
3: $K \leftarrow H(m, c)$
4: **return** (c, K)

Algorithm 16 $\texttt{Decap}^{\not\perp}(\texttt{sk}, c)$

1: $\hat{m} \leftarrow \texttt{D}(\texttt{sk}, c)$
2: **if** $\hat{m} \in \mathcal{M}_{\texttt{PKE}}$ **then**
3:      $K \leftarrow H(\hat{m}, c)$
4:      **return** K
5: **end if**
6: **return** $H(s, c)$

Figure 8: $U^{\not\perp}$ routines

The security of $\texttt{KEM}^{\not\perp}$, similar to the security of $\texttt{KEM}^\perp$, derives from the indistinguishability between pseudorandom $K \leftarrow H(m, c)$ and truly random $K \xleftarrow{\$} \mathcal{K}_{\texttt{KEM}}$ under the random oracle model, unless the $\texttt{KEM}$ adversary somehow queries H on (m^*, c^*), in which case an adversary against the underlying $\texttt{PKE}$ with access to a $\texttt{PCO}$ can detect this special query use it to win the $\texttt{OW-PCA}$ game. In addition, because $\texttt{Decap}^{\not\perp}$ will always return random-looking values regardless of the validity of the input ciphertext, the decapsulation oracle can be simulated without needing a ciphertext validation oracle. The $\texttt{KEM}$ adversary will be able to distinguish between a true decapsulation oracle and a simulated decapsulation oracle, but only by querying H on (s, c) for some c. Since s is uniformly random, each query has a $|\mathcal{M}_{\texttt{PKE}}|^{-1}$ chance of hitting s, so across q_H hash queries to H, the probability of hitting s at least once is at most $\frac{q_H}{|\mathcal{M}_{\texttt{PKE}}|}$.

Theorem 2.4. *For every $\texttt{IND-CCA}$ adversary B against $\texttt{KEM}^{\not\perp}$ that makes q_H hash queries to H, there exists an $\texttt{OW-PCA}$ adversary A against the underlying $\texttt{PKE}$ such that:*

$$\texttt{Adv}(B) \leq \frac{q_H}{|\mathcal{M}_{PKE}|} + \texttt{Adv}(A)$$

2.1.3 KEM with explicit rejection from rigid PKE

If the input PKE is rigid (definition 1.5), then the PCO can be simulated with indistinguishability from the true PCO except for when decryption failure occurs. Given $\text{PKE} = (\text{KeyGen}, \text{E}, \text{D})$ and hash function $H : \mathcal{M}_{\text{PKE}} \to \mathcal{K}_{\text{KEM}}$, the $U_m^\perp$ transformation outputs a $\text{KEM}_m^\perp = (\text{KeyGen}, \text{Encap}_m^\perp, \text{Decap}_m^\perp)$ where the key generation routine remains unchanged. The encapsulation and decapsulation routines are described in figure 9

Algorithm 17 $\text{Encap}_m^\perp(\text{pk})$

1: $m \xleftarrow{\$} \mathcal{M}_{\text{PKE}}$
2: $c \leftarrow \text{E}(\text{pk}, m)$
3: $K \leftarrow H(m)$
4: **return** (c, K)

Algorithm 18 $\text{Decap}_m^\perp(\text{sk}, c)$

1: $\hat{m} \leftarrow \text{D}(\text{sk}, c)$
2: **if** $\hat{m} \in \mathcal{M}_{\text{PKE}}$ **then**
3:      **return** $H(m)$
4: **end if**
5: **return** $\perp$

Figure 9: $\text{KEM}_m^\perp$ routines

Under the random oracle model, no adversary B can distinguish between the pseudorandom $K_0 \leftarrow H(m)$ and the truly random $K_1 \xleftarrow{\$} \mathcal{K}_{\text{KEM}}$ unless B queries H on m^*. Because rigidity means $\text{E}(\text{pk}, m^*) = c^*$ is equivalent to $\text{D}(\text{sk}, c^*) = m^*$, if H is queried on m^*, a second adversary A against the PKE can detect the special query and use it to win the OW-VA game.

Theorem 2.5. *For every IND-CCA2 adversary B against $\text{KEM}_m^\perp$, there exists an OW-VA adversary A against the underlying PKE such that*

$$\textit{Adv}(B) \leq \textit{Adv}(A) + \delta \cdot (q_H + q_D)$$

2.1.4 KEM with implicit rejection from rigid PKE

Figure 10: $U_m^{\not\perp}$ routines

2.2 Applications

Kyber's round 3 submission [ABD+19] and ML-KEM [KE23] uses *de-randomization*, *re-encryption*, and some variation of the U transformation to convert the IND-CPA PKE into an IND-CCA2 KEM. Here we recall the IND-CPA PKE first, then describe and discuss their respective KEM construction.

In the PKE, $R_q = \mathbb{Z}_q[x]/\langle x^n + 1 \rangle$ denotes the quotient ring, $\mathcal{X}_\eta$ denotes the centered binomial distribution that spans $\{-\eta, \eta + 1, \ldots, \eta - 1, \eta\}$. The plaintext space is the set of all polynomials in R_q with coefficients $0, 1$. k denotes the dimensions of the secret. The full set of parameters is listed in table 4

Security level	n	q	k	η_1	η_2
Kyber512	256	3329	2	3	2
Kyber768	256	3329	3	2	2
Kyber1024	256	3329	4	3	2

Table 4: Kyber/ML-KEM parameter set

The IND-CPA PKE(KeyGen, E, D) is identical between round 3 Kyber and ML-KEM. The details of the PKE routines are described in figure 11

Algorithm 19 KeyGen$_{\text{PKE}}$

1: $A \xleftarrow{\$} R_q^{k \times k}, \mathbf{s} \xleftarrow{\$} \mathcal{X}_{\eta_1}^k$
2: $\mathbf{t} \leftarrow A \cdot \mathbf{s}$
3: $\text{pk} \leftarrow (A, \mathbf{t}), \text{sk} \leftarrow \mathbf{s}$
4: **return** (pk, sk)

Algorithm 20 $\text{E}_{\text{PKE}}(\text{pk}, m)$

1: $(A, \mathbf{t}) \leftarrow \text{pk}$
2: $\mathbf{r}_1 \xleftarrow{\$} \mathcal{X}_{\eta_1}^k$
3: $\mathbf{e}_1 \xleftarrow{\$} \mathcal{X}_{\eta_2}^k, e_2 \xleftarrow{\$} \mathcal{X}_{\eta_2}$
4: $\mathbf{c}_1 \leftarrow A^\mathsf{T} \cdot \mathbf{r}_1 + \mathbf{e}_1$
5: $c_2 \leftarrow \mathbf{t}^\mathsf{T} \cdot \mathbf{r}_1 + e_2 + m \cdot \left\lceil \frac{q}{2} \right\rfloor$
6: **return** $(\mathbf{c}_1, c_2)$

Algorithm 21 $\text{D}_{\text{PKE}}(\text{sk}, c)$

1: $(\mathbf{c}_1, c_2) \leftarrow c$
2: $\mathbf{s} \leftarrow \text{sk}$
3: $\hat{m} = c_2 - \mathbf{s}^\mathsf{T} \cdot \mathbf{c}_1$
4: $\hat{m} \leftarrow \text{Round}(\hat{m})$
5: **return** $\hat{m}$

Figure 11: PKE routines

For constructing KEM from PKE, Kyber uses the $U^{\not\perp}$ transformation while ML-KEM uses the $U_m^{\not\perp}$ transformation. In implementation, however, there are two deviations from algorithm listed in figure 8 and 10.

The first deviation is the inclusion of the public key pk when deriving the coin. Instead of deriving coin from the plaintext only $r \leftarrow G(m)$, Kyber and ML-KEM derive the coin from both the plaintext and (a hash of) the public key $(K, r) \leftarrow G(m, H(\text{pk}))$. The main rationale [BDK$^+$18] behind this design choice is to prevent adversaries from precomputing m values that have a high likelihood of triggering decryption failure, which can reveal information about the secret key.

Recall from the decryption routine D_{PKE}:

$$\hat{m} = \mathbf{s}^\mathsf{T} \cdot \mathbf{c}_1 - c_2$$
$$= \mathbf{s}^\mathsf{T} \cdot (A^\mathsf{T} \cdot \mathbf{r}_1 + \mathbf{e}_1) - (\mathbf{t}^\mathsf{T} \cdot \mathbf{r}_1 + e_2 + m \cdot \left\lceil \frac{q}{2} \right\rfloor)$$
$$= \mathbf{s}^\mathsf{T} \cdot (A^\mathsf{T} \cdot \mathbf{r}_1 + \mathbf{e}_1) - ((A \cdot \mathbf{s} + \mathbf{e})^\mathsf{T} \cdot \mathbf{r}_1 + e_2 + m \cdot \left\lceil \frac{q}{2} \right\rfloor)$$
$$= \mathbf{s}^\mathsf{T} \cdot \mathbf{e}_1 - \mathbf{e}^\mathsf{T} \cdot \mathbf{r}_1 + e_2 + m \left\lceil \frac{q}{2} \right\rfloor$$

A decryption failure occurs if the noise term $\mathbf{s}^\mathsf{T} \cdot \mathbf{e}_1 - \mathbf{e}^\mathsf{T} \cdot \mathbf{r}_1 + e_2$ has large coefficient, which roughly translates to $\mathbf{s}^\mathsf{T} \cdot \mathbf{e}_1$ having large coefficients. If the coin $r \leftarrow G(m)$ is entirely derived from the adversary, then the adversary can compute a large number of $\mathbf{r}_1, \mathbf{e}_1, e_2$ offline, then observe which set of $\mathbf{r}_1, \mathbf{e}_1, e_2$ trigger decryption failure, which then can be used to infer information about $\mathbf{s}$ and $\mathbf{e}$. On the other hand, if the coin is derived from both the message and the public key, then the adversary needs to compute a large set of value per keypair, which pushes this attack into infeasibility even for adversaries with large quantum computers.

A second deviation is that instead of deriving the shared secret from the plaintext and the ciphertext $K \leftarrow H(m, c)$, Kyber derives the shared secret from the plaintext and *a hash of the ciphertext*. Since ML-KEM uses the $U_m^{\not\perp}$ transformation, this deviation does not apply. [BDK$^+$18] cites benchmarking for non-incremental hash APIs to be the reason for this deviation, which I do not fully understand, but has likely less to do with the security of the scheme and more with verification and compliance.

The details of Kyber and ML-KEM's KEM construction are described in figures 12 and 13 respectively.

Algorithm 22 KeyGen$_{\text{KyberKEM}}$

1: $(\text{pk}_{\text{PKE}}, \text{sk}_{\text{PKE}}) \xleftarrow{\$} \text{KeyGen}_{\text{PKE}}()$
2: $z \xleftarrow{\$} R_{0,1}$
3: $h \leftarrow H(\text{pk}_{\text{PKE}})$
4: $\text{sk}_{\text{KEM}} \leftarrow (\text{sk}_{\text{PKE}}, z, h, \text{pk}_{\text{PKE}})$
5: $\text{pk}_{\text{KEM}} \leftarrow \text{pk}_{\text{PKE}}$
6: **return** $(\text{pk}_{\text{KEM}}, \text{sk}_{\text{KEM}})$

Algorithm 23 Encap$_{\text{KyberKEM}}(\text{pk})$

1: $m \xleftarrow{\$} R_2$
2: $(K, r) \leftarrow G(H(\text{pk}), m)$
3: $c \leftarrow \text{E}_{\text{PKE}}(\text{pk}, m; r)$
4: $K \leftarrow \text{KDF}(K, c)$
5: **return** (c, K)

Algorithm 24 Decap$_{\text{KyberKEM}}(\text{sk}, c)$

1: $(\text{sk}_{\text{PKE}}, z, h, \text{pk}_{\text{PKE}}) \leftarrow \text{sk}$
2: $\hat{m} \leftarrow \text{D}_{\text{PKE}}(\text{sk}_{\text{PKE}}, c)$
3: $(\hat{K}, \hat{r}) \leftarrow G(h, \hat{m})$
4: $\hat{c} \leftarrow \text{E}_{\text{PKE}}(\text{pk}, \hat{m}; \hat{r})$
5: **if** $\hat{c} = c$ **then**
6: $\quad$ **return** $\text{KDF}(\hat{K}, H(\hat{c}))$
7: **else**
8: $\quad$ **return** $\text{KDF}(z, H(\hat{c}))$
9: **end if**

Figure 12: Kyber's KEM routines

Algorithm 25 KeyGen$_{\text{ML-KEM}}$

1: $(\text{pk}_{\text{PKE}}, \text{sk}_{\text{PKE}}) \xleftarrow{\$} \text{KeyGen}_{\text{PKE}}()$
2: $z \xleftarrow{\$} R_2$
3: $h \leftarrow H(\text{PKE}_{\text{PKE}})$
4: $\text{sk}_{\text{KEM}} \leftarrow (\text{sk}_{\text{PKE}}, z, h, \text{pk}_{\text{PKE}})$
5: $\text{pk}_{\text{KEM}} \leftarrow \text{PKE}_{\text{PKE}}$
6: **return** $(\text{pk}_{\text{KEM}}, \text{sk}_{\text{KEM}})$

Algorithm 26 Encap$_{\text{ML-KEM}}(\text{pk}_{\text{KEM}})$

1: $m \xleftarrow{\$} R_2$
2: $(K, r) \leftarrow G(m, H(\text{pk}))$
3: $c \leftarrow \text{E}_{\text{PKE}}(\text{pk}, m; r)$
4: **return** (c, K)

Algorithm 27 Decap$_{\text{ML-KEM}}(\text{sk}, c)$

1: $(\text{sk}_{\text{PKE}}, z, h, \text{pk}_{\text{PKE}}) \leftarrow \text{sk}$
2: $\hat{m} \leftarrow \text{D}_{\text{PKE}}(\text{sk}_{\text{PKE}}, c)$
3: $(\hat{K}, \hat{r}) \leftarrow G(\hat{m}, h)$
4: $K' \leftarrow \text{KDF}(z, c)$
5: $\hat{c} \leftarrow \text{E}_{\text{PKE}}(\text{pk}_{\text{PKE}}, \hat{m}; \hat{r})$
6: **if** $\hat{c} = c$ **then**
7:     **return** $\hat{K}$
8: **else**
9:     **return** K'
10: **end if**

Figure 13: ML-KEM's KEM routines

3 Plaintext awareness and OAEP

Another successful strategy for constructing `IND-CCA2` secure encryption scheme is *Optimal Asymmetric Encryption Padding*[BR95]. OAEP achieves security against chosen ciphertext attacks by achieving a special property called *Plaintext Awareness*. We will introduce plaintext awareness in 3.1 and discuss its relationship with `IND-CCA` security, then describe OAEP and show that it is plaintext aware under the ROM.

3.1 Plaintext awareness

From a high level, plaintext awareness describes a public-key encryption scheme in which an adversary cannnot produce valid ciphertexts without "knowing" the corresponding plaintexts. The "knowledge of plaitnext" is modeled by a universal knowledge extractor K, which uses the public key `pk` and the random oracle tape $\mathcal{L}$ to recover the decryptin of some ciphertext c without using the secret key. Bellare noted in [BDPR98] that such a knowledge extractor can only exist within the ROM, or the one-way security of the encryption scheme is trivially broken.

Plaintext awareness also models the adversary's ability to eavesdrop using an "encryption oracle $\mathcal{O}^{\mathcal{E}}$", from which the adversary can query valid ciphertexts without leaving the corresponding hash queries in the random oracle. We rule out trivials wins in which an adversary samples a random ciphertexts to which nobody knows the answer by requiring the adversary's output to not be in the encryption oracle's tape.

We can now formally define the `PA` game. Let $\text{PKE}(\text{KeyGen}, \text{E}, \text{D})$ be a public key encryption scheme. Let H be a hash function. Let $\mathcal{O}^{\mathcal{E}}$ be the an encryption oracle. Let $\mathcal{L}^H$ and $\mathcal{L}^{\mathcal{E}}$ respectively denote the tape of the random oracle and the tape of the encryption oracle.

Algorithm 28 The `PA` game

1: $(\text{pk}, \text{sk}) \xleftarrow{\$} \text{KeyGen}()$
2: $c \xleftarrow{\$} A^{\mathcal{O}^H, \mathcal{O}^{\mathcal{E}}}(\text{pk})$
3: $\hat{m} \leftarrow K(\text{pk}, c, \mathcal{L}^H, \mathcal{L}^{\mathcal{E}})$
4: **return** $[\![c \notin \mathcal{L}^{\mathcal{E}} \wedge \hat{m} \neq D(\text{sk}, c)]\!]$

Figure 14: The `PA` game

`PA` adversary wins the game if c is not in $\mathcal{L}^{\mathcal{E}}$ and the knowledge extractor fails to recover the correct decryption.

Definition 3.1 (ϵ-Plaintext Awareness)**.** *A* `PKE` *is* ϵ-*Plaintext Aware if there exists a knowledge extractor* K *such that for all efficient* `PA` *adversary* A*, the probability of winning the game is at most* ϵ*.*

If a PKE is ϵ-plaintext aware with a negligible ϵ, then a decryption oracle can be simulated using the universal knowledge extractor with negligible difference in behaviors.

Theorem 3.1 ([BDPR98] theorem 4.2). *If PKE is ϵ-plaintext aware, then for every IND-CCA2 adversary B against PKE making at most q_D decryption queries, there exists an IND-CPA adversary A such that*

$$\mathtt{Adv}(B) \leq \mathtt{Adv}(A) + q_D \cdot \epsilon$$

3.2 Optimal Asymmetric Encryption Padding

References

[ABD⁺19] Roberto Avanzi, Joppe Bos, Léo Ducas, Eike Kiltz, Tancrède Lepoint, Vadim Lyubashevsky, John M Schanck, Peter Schwabe, Gregor Seiler, and Damien Stehlé. Crystals-kyber algorithm specifications and supporting documentation. *NIST PQC Round*, 2(4):1–43, 2019.

[BDK⁺18] Joppe Bos, Léo Ducas, Eike Kiltz, Tancrède Lepoint, Vadim Lyubashevsky, John M Schanck, Peter Schwabe, Gregor Seiler, and Damien Stehlé. Crystals-kyber: a cca-secure module-lattice-based kem. In *2018 IEEE European Symposium on Security and Privacy (EuroS&P)*, pages 353–367. IEEE, 2018.

[BDPR98] Mihir Bellare, Anand Desai, David Pointcheval, and Phillip Rogaway. Relations among notions of security for public-key encryption schemes. In *Advances in Cryptology—CRYPTO'98: 18th Annual International Cryptology Conference Santa Barbara, California, USA August 23–27, 1998 Proceedings 18*, pages 26–45. Springer, 1998.

[BR95] Mihir Bellare and Phillip Rogaway. Optimal asymmetric encryption. In *Advances in Cryptology—EUROCRYPT'94: Workshop on the Theory and Application of Cryptographic Techniques Perugia, Italy, May 9–12, 1994 Proceedings 13*, pages 92–111. Springer, 1995.

[FO99] Eiichiro Fujisaki and Tatsuaki Okamoto. Secure integration of asymmetric and symmetric encryption schemes. In *Annual international cryptology conference*, pages 537–554. Springer, 1999.

[HHK17] Dennis Hofheinz, Kathrin Hövelmanns, and Eike Kiltz. A modular analysis of the fujisaki-okamoto transformation. In *Theory of Cryptography Conference*, pages 341–371. Springer, 2017.

[KE23] NIST Module-Lattice-Based Key-Encapsulation. Mechanism standard. *NIST Post-Quantum Cryptography Standardization Process; NIST: Gaithersburg, MD, USA*, 2023.