

Faster generic IND-CCA2 secure KEM using “encrypt-then-MAC”

Anonymous Submission

Abstract. The modular Fujisaki-Okamoto (FO) transformation takes public-key encryption with weaker security and constructs a key encapsulation mechanism (KEM) with indistinguishability under adaptive chosen ciphertext attacks. While the modular FO transform enjoys tight security bound and quantum resistance, it also suffers from computational inefficiency and vulnerabilities to side-channel attacks due to using de-randomization and re-encryption for providing ciphertext integrity. In this work, we propose an alternative KEM construction that achieves ciphertext integrity using a message authentication code (MAC) and instantiate a concrete instance using Kyber. Our experimental results showed that where the encryption routine incurs heavy computational cost, replacing re-encryption with MAC provides substantial performance improvements at comparable security level.

Keywords: Key encapsulation mechanism, post-quantum cryptography, lattice cryptography, Fujisaki-Okamoto transformation

1 Introduction

A key encapsulation mechanism (KEM) [Sho01] is a cryptographic primitive that allows two parties to establish a shared secret over an insecure channel. The desired security standard for a KEM is called indistinguishability under adaptive chosen-ciphertext attack (IND-CCA2). Intuitively speaking, IND-CCA2 security requires that no efficient adversary can distinguish a pseudorandom shared secret from truly random noise, even with access to a decapsulation oracle throughout the attack. However, building a provably IND-CCA2 KEM is tremendously difficult. Early attempts without formal proof, such as RSA encryption defined in PKCS#1 v1.5 [Kal98], were later shown to be vulnerable to practical chosen-ciphertext attack [Ble98]. In recent decades, the most viable approach has been to start with cryptographic primitives possessing weaker security properties, such as a public-key encryption (PKE) scheme with one-way security under chosen-plaintext attack (OW-CPA), then add steps to ensure ciphertext non-malleability [?]. Some of the earliest proposals for generic IND-CCA2 secure constructions include OAEP [?], Fujisaki-Okamoto transformation [FO99][FO13], REACT [OP01b], and GEM [CHJ⁺02].

On the other hand, chosen-ciphertext security is a solve problem in symmetric cryptography. It is well understood that, by combining an IND-CPA secure symmetric encryption scheme with an existentially unforgeable message authentication code (MAC) in a pattern called “encrypt-then-MAC” [Kra01], one can build an authenticated encryption scheme [BN00] that achieves IND-CCA2 security. While this technique cannot be directly applied in the context of public-key cryptography due to the lack of a shared symmetric key between the two communicating parties, the concept of authenticating ciphertext using a MAC still has strong merits. Abdalla, Rogaway, and Bellare proposed DHIES (also known as “Hashed ElGamal”)[ABR99][ABR01], a hybrid public-key encryption (HPKE) scheme whose IND-CCA2 security reduces to the Gap Diffie-Hellman assumption [OP01a] under the random oracle model. The technique behind DHIES is to derive both the shared secret and a symmetric MAC key by hashing a random PKE plaintext, encrypt the PKE

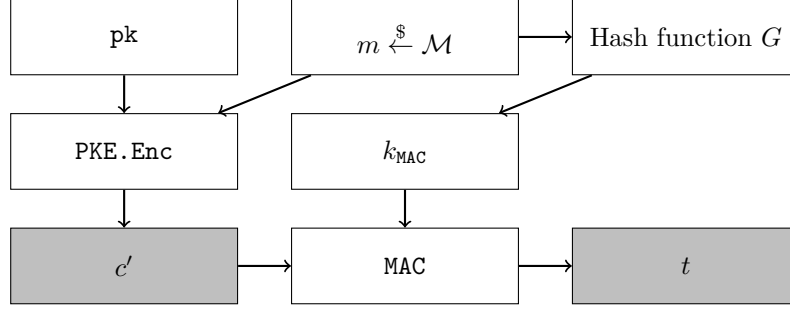


Figure 1: Combining PKE with MAC using “encrypt-then-MAC” to encapsulate a shared secret. The returned values are colored grey

44 plaintext, then authenticate the PKE ciphertext using the previously derived MAC key.
 45 Where the Gap Diffie-Hellman assumption holds and the MAC is existentially unforgeable,
 46 no efficient adversary can recover the decryption of an unknown ciphertext even with
 47 access to a decryption oracle because it cannot produce a valid tag for such unknown
 48 ciphertext.

49 1.1 Our contribution

50 In this paper, we propose a generic IND-CCA2 secure KEM using techniques adapted
 51 from DHIES. We call our proposed scheme the “encrypt-then-MAC” transformation
 52 due to the conceptual similarity to the namesake symmetric encryption technique for
 53 achieving authenticated encryption. A summary of the data flow in the “encrypt-then-MAC”
 54 transformation can be found in figure 1.

55 Our scheme mainly differs from DHIES in its versatility. Whereas the IND-CCA2
 56 security of DHIES reduces specifically to the Gap Diffie-Hellman assumption, the chosen-
 57 ciphertext security of the “encrypt-then-MAC” KEM reduces more generally to the
 58 OW-PCA security [OP01b] of the input scheme (we will show in section 3 that the Gap
 59 Diffie-Hellman assumption implies a special case of OW-PCA security). In addition,
 60 we propose that because each call to encapsulation samples a fresh PKE plaintext, the
 61 “encrypt-then-MAC” KEM can be instantiated with one-time secure MAC such as Poly1305
 62 (Abdalla et al originally proposed to use HMAC and CBC-MAC, which are many-time
 63 secure MAC but less efficient than one-time MAC, see section 4.1), which greatly improves
 64 the computational efficiency of the scheme.

65 1.1.1 Performance improvements

66 In section 4, we instantiate an instance of “encrypt-then-MAC” KEM using the PKE
 67 subroutines defined in ML-KEM [oST24]. For clarity’s sake we call this instantiation
 68 ML-KEM⁺. Compared with the full KEM routine in ML-KEM, which is based on some
 69 variant of the modular Fujisaki-Okamoto KEM transformation, ML-KEM⁺ demonstrates
 70 enormous runtime efficiency gain in decapsulation with only minimal performance penalty
 71 in encapsulation runtime and ciphertext size. On an AMD EPYC 9R14 CPU, ML-KEM⁺
 72 (combined with Poly1305) achieves between 72%-80% reduction of CPU cycles needed for
 73 decapsulation while only incurring 2%-7% increase of CPU cycles needed for encapsulation.
 74 See table 1 for a summary of CPU cycles needed for encapsulation and decapsulation in
 75 each security level.

Table 1: ML-KEM⁺ is instantiated with Poly1305

	ML-KEM 512	ML-KEM ⁺ 512	ML-KEM 768	ML-KEM ⁺ 768	ML-KEM 1024	ML-KEM ⁺ 1024
Encap (ccl/tick)	91467	93157 (+1.8%)	136405	146405 (+7.3%)	199185	205763 (+3.3%)
Decap (ccl/tick)	121185	33733 (-72.2%)	186445	43315 (-76.8%)	246245	51375 (-79.1%)
CT size (bytes)	768	784 (+2.1%)	1088	1104 (+1.5%)	1568	1584 (+1.0%)

Since the decapsulation efficiency comes at the cost of increasing encryption runtime and ciphertext size, we also measured the round trip time of key exchange protocols with various modes of authentication. When compared to ML-KEM, ML-KEM⁺ achieves 24%-28% reduction of round trip time in unauthenticated key exchange (KE), 29%-35% in unilaterally authenticated key exchange (UAKE), and 35%-48% reduction in mutually authenticated key exchange (AKE). See table 2 for a summary of round trip times.

Table 2: Key exchange round-trip times

	ML-KEM 512	ML-KEM ⁺ 512	ML-KEM 768	ML-KEM ⁺ 768	ML-KEM 1024	ML-KEM ⁺ 1024
KE RTT (μ s)	92	70 (-23.9%)	135	99 (-26.7%)	193	138 (-28.5%)
UAKE RTT (μ s)	145	103 (-29.0%)	215	144 (-33.0%)	310	202 (-34.8%)
AKE RTT (μ s)	220	133 (-39.5%)	294	190 (-35.4%)	512	266 (-48.0%)

1.2 Related works

Optimal Asymmetric Encryption Padding (OAEP) [BR94][BDPR98] is a generic chosen-ciphertext secure PKE. Under the random oracle model, the chosen-ciphertext security of the OAEP encryption scheme reduces to the one-wayness of the input trapdoor permutation. Although it was discovered that there exist trapdoor permutations with which the OAEP encryption scheme does not achieve full IND-CCA2 security [Sho02], Fujisaki et al later proved that the OAEP is IND-CCA2 secure when combined with the RSA trapdoor permutation [FOPS01][RSA78]. RSA-OAEP was standardized in PKCS#1 v2 [MKJR16] and is currently the most recommended method of RSA-based encryption scheme. Unfortunately, OAEP’s requirement for a trapdoor permutation is immensely difficult to satisfy, and no other practical instantiation saw widespread adoption to this day.

The **Fujisaki-Okamoto transformation** [FO99][FO13] is another generic chosen-ciphertext secure transformation. While Fujisaki and Okamoto originally proposed a hybrid public-key encryption scheme whose IND-CCA2 security reduces non-tightly to the OW-CPA security of the input PKE and the IND-CPA security of the symmetric encryption scheme. Later works [Den03][HHK17a][DNR04][HHM22][BP18] tightened the security reduction, accounted for imperfect correctness, adapted the original proposal to build KEM, and proved its security in the quantum random oracle model (QROM).

The modular Fujisaki-Okamoto KEM transformation is remarkably successful because of the simplicity of its construction, the tightness of the security bound, and the proven (though non-tight) security against quantum adversary. It was adopted by many submissions to NIST’s post-quantum cryptography competition, including Kyber [BDK⁺18], Saber [DKRV18], FrodoKEM [BCD⁺16], and classic McEliece [ABC⁺20] among others.

However, the Fujisaki-Okamoto transformation is not perfect. Among its shortfalls are two problems caused by its use of *de-randomization* and *re-encryption* for achieving chosen-ciphertext security.

- **Computational inefficiency.** The decapsulation routine needs to re-encrypt the decryption to ensure ciphertext has not been tempered with. For input PKE whose encryption routine carries substantial computational cost, such as most lattice-based cryptosystems, re-encryption slows down decapsulation significantly.
- **Side-channel vulnerability.** Re-encryption also introduces side-channels that can leak information about the decrypted PKE plaintext. As demonstrated in [UXT+22][RRCB19][TUX+23], these side-channels can be converted into efficient plaintext-checking attacks that can fully recover the secret key in fewer than 1000 traces.

2 Preliminaries

2.1 Public-key encryption scheme

Syntax. A public-key encryption scheme $\text{PKE}(\text{KeyGen}, \text{Enc}, \text{Dec})$ is a collection of three routines defined over some plaintext space \mathcal{M} and some ciphertext space \mathcal{C} . $(\text{pk}, \text{sk}) \xleftarrow{\$} \text{KeyGen}()$ is a randomized routine that returns a keypair. The encryption routine $\text{Enc} : (\text{pk}, m) \mapsto c$ encrypts the input plaintext under the input public key. The decryption routine $\text{Dec} : (\text{sk}, c) \mapsto m$ decrypts the input ciphertext under the input secret key. Where the encryption routine is randomized, we denote the randomness by $r \in \mathcal{R}$, where \mathcal{R} is called the coin space. The decryption routine is assumed to always be deterministic. Some decryption routines can detect malformed ciphertext and output the rejection symbol \perp accordingly.

Correctness. Following the definition in [DNR04] and [HHK17b], a PKE is δ -correct if:

$$E \left[\max_{m \in \mathcal{M}} P \left[\text{Dec}(\text{sk}, c) \neq m \mid c \xleftarrow{\$} \text{Enc}(\text{pk}, m) \right] \right] \leq \delta$$

Where the expectation is taken with respect to the probability distribution of all possible keypairs $(\text{pk}, \text{sk}) \xleftarrow{\$} \text{PKE}.\text{KeyGen}()$. For many lattice-based cryptosystems, including ML-KEM, decryption failures could leak information about the secret key, although the probability of a decryption failure is low enough that classical adversaries cannot exploit decryption failure more than they can defeat the underlying lattice problem. On the other hand, a quantum adversary may be able to exploit decryption failure in reasonable runtime by efficiently searching through all possible inputs using Grover’s search algorithm. For that, ML-KEM made slight modifications in its KEM construction to prevent quantum adversary from precomputing large lookup table. We refer readers to [ABD⁺19] and [BDK⁺18] for the details.

Security. We discuss the security of a PKE using the sequence of games described in [Sho04]. Specifically, we first define the OW-ATK as they pertain to a public key encryption scheme. In later section we will define the IND-CCA game as it pertains to a key encapsulation mechanism.

The OW-ATK game	
1: $(\mathbf{pk}, \mathbf{sk}) \xleftarrow{\$} \text{KeyGen}(1^\lambda)$	
2: $m^* \xleftarrow{\$} \mathcal{M}$	
3: $c^* \xleftarrow{\$} \text{Enc}(\mathbf{pk}, m^*)$	
4: $\hat{m} \xleftarrow{\$} \mathcal{A}^{\mathcal{O}_{\text{ATK}}}(1^\lambda, \mathbf{pk}, c^*)$	$\text{PCO}(m \in \mathcal{M}, c \in \mathcal{C})$
5: return $\llbracket m^* = \hat{m} \rrbracket$	1: return $\llbracket \text{Dec}(\mathbf{sk}, c) = m \rrbracket$

Figure 2: One-way security game of PKE (left) and plaintext-checking oracle (right)

145 In the OW-ATK game (see figure 2), an adversary's goal is to recover the decryption of a
 146 randomly generated ciphertext. A challenger randomly samples a keypair and a challenge
 147 plaintext m^* , encrypts the challenge plaintext $c^* \xleftarrow{\$} \text{Enc}(\mathbf{pk}, m^*)$, then gives \mathbf{pk} and c^*
 148 to the adversary A . The adversary A , with access to some oracle \mathcal{O}_{ATK} , outputs a guess
 149 decryption \hat{m} . A wins the game if its guess \hat{m} is equal to the challenge plaintext m^* . The
 150 *advantage* $\text{Adv}_{\text{OW-ATK}}$ of an adversary in this game is the probability that it wins the game:

$$\text{Adv}_{\text{OW-ATK}}(A) = P \left[A(\mathbf{pk}, c^*) = m^* \mid (\mathbf{pk}, \mathbf{sk}) \xleftarrow{\$} \text{KeyGen}(); m^* \xleftarrow{\$} \mathcal{M}; c^* \xleftarrow{\$} \text{Enc}(\mathbf{pk}, m^*) \right]$$

151 The capabilities of the oracle \mathcal{O}_{ATK} depends on the choice of security goal ATK. Particu-
 152 larly relevant to our result is security against plaintext-checking attack (PCA), for which
 153 the adversary has access to a plaintext-checking oracle (PCO) (see figure 2). A PCO takes
 154 as input a plaintext-ciphertext pair (m, c) and returns **True** if m is the decryption of c or
 155 **False** otherwise.

156 2.2 Key encapsulation mechanism (KEM)

157 A key encapsulation mechanism is a collection of three routines (**KeyGen**, **Encap**, **Decap**)
 158 defined over some ciphertext space \mathcal{C} and some key space \mathcal{K} . The key generation routine
 159 takes the security parameter 1^λ and outputs a keypair $(\mathbf{pk}, \mathbf{sk}) \xleftarrow{\$} \text{KeyGen}(1^\lambda)$. **Encap**(\mathbf{pk})
 160 is a probabilistic routine that takes a public key \mathbf{pk} and outputs a pair of values (c, K)
 161 where $c \in \mathcal{C}$ is the ciphertext (also called encapsulation) and $K \in \mathcal{K}$ is the shared secret
 162 (also called session key). **Decap**(\mathbf{sk}, c) is a deterministic routine that takes the secret key
 163 \mathbf{sk} and the encapsulation c and returns the shared secret K if the ciphertext is valid. Some
 164 KEM constructions use explicit rejection, where if c is invalid then **Decap** will return a
 165 rejection symbol \perp ; other KEM constructions use implicit rejection, where if c is invalid
 166 then **Decap** will return a fake session key that depends on the ciphertext and some other
 167 secret values.

168 The IND-CCA security of a KEM is defined by an adversarial game in which an
 169 adversary's goal is to distinguish pseudorandom shared secret (generated by running the
 170 **Encap** routine) and a truly random value.

KEM-IND-CCA2 game	
1: $(\mathbf{pk}, \mathbf{sk}) \xleftarrow{\$} \text{KeyGen}(1^\lambda)$	
2: $(c^*, K_0) \xleftarrow{\$} \text{Encap}(\mathbf{pk})$	
3: $K_1 \xleftarrow{\$} \mathcal{K}$	
4: $b \xleftarrow{\$} \{0, 1\}$	
5: $\hat{b} \xleftarrow{\$} A^{\mathcal{O}_{\text{Decap}}}(1^\lambda, \mathbf{pk}, c^*, K_b)$	$\mathcal{O}_{\text{Decap}}(c)$
6: return $\llbracket \hat{b} = b \rrbracket$	1: return $\text{Decap}(\mathbf{sk}, c)$

Figure 3: IND-CCA2 game for KEM (left) and decapsulation oracle (right)

171 The decapsulation oracle $\mathcal{O}^{\text{Decap}}$ takes a ciphertext c and returns the output of the
 172 **Decap** routine using the secret key. The advantage $\epsilon_{\text{IND-CCA}}$ of an IND-CCA adversary
 173 $\mathcal{A}_{\text{IND-CCA}}$ is defined by

$$\text{Adv}_{\text{IND-CCA}}(A) = \left| P[A^{\mathcal{O}_{\text{Decap}}}(a^\lambda, \mathbf{pk}, c^*, K_b) = b] - \frac{1}{2} \right|$$

174 2.3 Message authentication code (MAC)

175 A message authentication code **MAC** is a collection of routines (**Sign**, **Verify**) defined over
 176 some key space \mathcal{K} , some message space \mathcal{M} , and some tag space \mathcal{T} . The signing routine
 177 **Sign**(k, m) takes the secret key $k \in \mathcal{K}$ and some message, and outputs a tag t . The
 178 verification routine **Verify**(k, m, t) takes the triplet of secret key, message, and tag, and
 179 outputs 1 if the message-tag pair is valid under the secret key, or 0 otherwise. Many MAC
 180 constructions are deterministic. For these constructions it is simpler to denote the signing
 181 routine by $t \leftarrow \text{MAC}(k, m)$ and perform verification using a simple comparison.

182 The security of a MAC is defined in an adversarial game in which an adversary, with
 183 access to some signing oracle $\mathcal{O}_{\text{Sign}}(m)$, tries to forge a new valid message-tag pair that
 184 has never been queried before. The existential unforgeability under chosen message attack
 185 (EUF-CMA) game is shown below:

EUF-CMA game	
1: $k^* \xleftarrow{\$} \mathcal{K}$	
2: $(\hat{m}, \hat{t}) \xleftarrow{\$} \mathcal{A}^{\mathcal{O}_{\text{Sign}}}()$	
3: return $\llbracket \text{Verify}(k^*, \hat{m}, \hat{t}) \wedge (\hat{m}, \hat{t}) \notin \mathcal{O}_{\text{Sign}} \rrbracket$	

Figure 4: The existential forgery game

186 The advantage $\text{Adv}_{\text{EUF-CMA}}$ of the existential forgery adversary is the probability that it
 187 wins the EUF-CMA game.

188 3 The “encrypt-then-MAC” transformation

189 Let \mathcal{B}^* denote the set of finite bit strings. Let $\text{PKE}(\text{KeyGen}, \text{Enc}, \text{Dec})$ be a public-key
 190 encryption scheme defined over message space \mathcal{M} and ciphertext space \mathcal{C} . Let $\text{MAC} : \mathcal{K}_{\text{MAC}} \times \mathcal{B}^* \rightarrow \mathcal{T}$
 191 be a deterministic message authentication code that takes a key $k \in \mathcal{K}_{\text{MAC}}$,
 192 some message $m \in \mathcal{B}^*$, and outputs a digest $t \in \mathcal{T}$. Let $G : \mathcal{M} \rightarrow \mathcal{K}_{\text{MAC}}$ be a hash

193 function that maps from PKE's plaintext space to MAC's key space. Let $H : \mathcal{B}^* \rightarrow \mathcal{K}_{\text{KEM}}$
 194 be a hash function that maps bit strings into the set of possible shared secrets. The
 195 "encrypt-then-MAC" transformation $\text{EtM}[\text{PKE}, \text{MAC}, G, H]$ constructs a key encapsulation
 196 mechanism $\text{KEM}_{\text{EtM}}(\text{KeyGen}_{\text{KEM}}, \text{Encap}, \text{Decap})$, whose routines are described in figure 5.

$\text{KEM}_{\text{EtM}}.\text{KeyGen}()$	$\text{KEM}_{\text{EtM}}.\text{Decap}(\text{sk}, c)$
1: $(\text{pk}, \text{sk}') \xleftarrow{\$} \text{PKE}.\text{KeyGen}()$ 2: $z \xleftarrow{\$} \mathcal{M}$ 3: $\text{sk} \leftarrow \text{sk}' \ z$ 4: return (pk, sk)	1: $(c', t) \leftarrow c$ 2: $(\text{sk}', z) \leftarrow \text{sk}$ 3: $\hat{m} \leftarrow \text{PKE}.\text{Dec}(\text{sk}', c')$ 4: $\hat{k} \leftarrow G(\hat{m})$ 5: if $\text{MAC}(\hat{k}, c') \neq t$ then 6: $K \leftarrow H(z, c')$ 7: else 8: $K \leftarrow H(\hat{m}, c')$ 9: end if 10: return K
$\text{KEM}_{\text{EtM}}.\text{Encap}(\text{pk})$	
1: $m \xleftarrow{\$} \mathcal{M}$ 2: $k \leftarrow G(m)$ 3: $c' \xleftarrow{\$} \text{PKE}.\text{Enc}(\text{pk}, m)$ 4: $t \leftarrow \text{MAC}(k, c')$ 5: $K \leftarrow H(m, c')$ 6: $c \leftarrow c' \ t$ 7: return (c, K)	

Figure 5: KEM_{EtM} routines

197 The key generation routine of KEM_{EtM} is largely identical to that of the PKE, only a
 198 secret value z is sampled as the implicit rejection symbol. In the encapsulation routine,
 199 a MAC key is derived from the randomly sampled plaintext $k \leftarrow G(m)$, then used
 200 to sign the unauthenticated ciphertext c' . Because the encryption routine might be
 201 randomized, the session key is derived from both the message and the ciphertext. Finally,
 202 the unauthenticated ciphertext c' and the tag t combine into the authenticated ciphertext
 203 c that would be transmitted to the peer. In the decapsulation routine, the decryption \hat{m}
 204 of the unauthenticated ciphertext is used to re-derive the MAC key \hat{k} , which is then used
 205 to re-compute the tag \hat{t} . The ciphertext is considered valid if and only if the recomputed
 206 tag is identical to the input tag.

207 For an adversary A to produce a valid tag t for some unauthenticated ciphertext
 208 c' under the symmetric key $k \leftarrow G(\text{Dec}(\text{sk}', c'))$ implies that A must either know the
 209 symmetric key k or produce a forgery. Under the random oracle model, A also cannot
 210 know k without knowing its preimage $\text{Dec}(\text{sk}', c')$, so A must either have produced c'
 211 honestly, or have broken the one-way security of PKE. This means that the decapsulation
 212 oracle will not give out information on decryptions that the adversary does not already
 213 know.

$\text{PCO}(m, c)$
1: $k \leftarrow G(m)$ 2: $t \leftarrow \text{MAC}(k, c)$ 3: return $\llbracket \mathcal{O}^{\text{Decap}}((c, t)) = H(m, c) \rrbracket$

Figure 6: Every decapsulation oracle can be converted into a plaintext-checking oracle

However, a decapsulation oracle can still give out some information: for a known plaintext m , all possible encryptions $c' \xleftarrow{\$} \text{Enc}(\text{pk}, m)$ can be correctly signed, while ciphertexts that don't decrypt back to m cannot be correctly signed. This means that a decapsulation oracle can be converted into a plaintext-checking oracle (see figure 6), so every chosen-ciphertext attack against the KEM can be converted into a plaintext-checking attack against the underlying PKE.

On the other hand, if the underlying PKE is one-way secure against plaintext-checking attack that makes q plaintext-checking queries, then “encrypt-then-MAC” KEM is semantically secure under chosen ciphertext attacks making the same number of decapsulation queries:

Theorem 1. *For every IND-CCA2 adversary A against KEM_{EtM} that makes q decapsulation queries, there exists an OW-PCA adversary B who makes at least q plaintext-checking queries against the underlying PKE, and an one-time existential forgery adversary C against the underlying MAC such that*

$$\text{Adv}_{\text{IND-CCA2}}(A) \leq q \cdot \text{Adv}_{\text{OT-MAC}}(C) + 2 \cdot \text{Adv}_{\text{OW-PCA}}(B)$$

Theorem 1 naturally flows into an equivalence relationship between the security of the KEM and the security of the PKE:

Lemma 1. *KEM_{EtM} is IND-CCA2 secure if and only if the input PKE is OW-PCA secure*

3.1 Proof of theorem 1

We will prove theorem 1 using a sequence of game. A summary of the the sequence of games can be found in figure 7 and 8. From a high level we made three incremental modifications to the IND-CCA2 game for KEM_{EtM} : replace true decapsulation with simulated decapsulation, replace the pseudorandom MAC key $k^* \leftarrow G(m^*)$ with a truly random MAC key $k^* \xleftarrow{\$} \mathcal{K}_{\text{MAC}}$, and finally replace pseudorandom shared secret $K_0 \leftarrow H(m^*, c')$ with a truly random shared secret $K_0 \xleftarrow{\$} \mathcal{K}_{\text{KEM}}$. A OW-PCA adversary can then simulate the modified IND-CCA2 game for the KEM adversary, and the advantage of the OW-PCA adversary is associated with the probability of certain behaviors of the KEM adversary.

Proof. *Game 0* is the standard IND-CCA2 game for KEMs. The decapsulation oracle $\mathcal{O}^{\text{Decap}}$ executes the decapsulation routine using the challenge keypair and return the results faithfully. The queries made to the hash oracles $\mathcal{O}^G, \mathcal{O}^H$ are recorded to their respective tapes $\mathcal{L}^G, \mathcal{L}^H$.

Game 1 is identical to game 0 except that the true decapsulation oracle $\mathcal{O}^{\text{Decap}}$ is replaced with a simulated oracle $\mathcal{O}_1^{\text{Decap}}$. Instead of directly decrypting c' as in the decapsulation routine, the simulated oracle searches through the tape \mathcal{L}^G to find a matching query (\tilde{m}, \tilde{k}) such that \tilde{m} is the decryption of c' . The simulated oracle then uses \tilde{k} to validate the tag t against c' .

If the simulated oracle accepts the queried ciphertext as valid, then there is a matching query that also validates the tag, which means that the queried ciphertext is honestly generated. Therefore, the true oracle must also accept the queried ciphertext. On the other hand, if the true oracle rejects the queried ciphertext (and output the implicit rejection $H(z, c')$), then the tag is simply invalid under the MAC key $k = G(\text{Dec}(\text{sk}', c'))$. Therefore, there could not have been a matching query that also validates the tag, and the simulated oracle must also rejects the queried ciphertext.

This means that from the adversary A 's perspective, game 1 and game 0 differ only when the true oracle accepts while the simulated oracle rejects, which means that t is a valid tag for c' under $k = G(\text{Dec}(\text{sk}', c'))$, but k has never been queried. Under the random

IND-CCA2 game for KEM_{EtM}	$\mathcal{O}^{\text{Decap}}(c)$
1: $(\text{pk}, \text{sk}) \xleftarrow{\$} \text{KEM}_{\text{EtM}}.\text{KeyGen}()$ 2: $m^* \xleftarrow{\$} \mathcal{M}$ 3: $c' \xleftarrow{\$} \text{PKE}.\text{Enc}(\text{pk}, m^*)$ 4: $k^* \leftarrow G(m^*)$ \triangleright Game 0-1 5: $k^* \xleftarrow{\$} \mathcal{K}_{\text{MAC}}$ \triangleright Game 2-3 6: $t \leftarrow \text{MAC}(k^*, c')$ 7: $c^* \leftarrow c' t$ 8: $K_0 \leftarrow H(m^*, c')$ \triangleright Game 0-2 9: $K_0 \xleftarrow{\$} \mathcal{K}_{\text{KEM}}$ \triangleright Game 3 10: $K_1 \xleftarrow{\$} \mathcal{K}_{\text{KEM}}$ 11: $b \xleftarrow{\$} \{0, 1\}$ 12: $\hat{b} \leftarrow A^{\mathcal{O}^{\text{Decap}}}(\text{pk}, c^*, K_b)$ \triangleright Game 0 13: $\hat{b} \leftarrow A^{\mathcal{O}_1^{\text{Decap}}}(\text{pk}, c^*, K_b)$ \triangleright Game 1-3 14: return $[\hat{b} = b]$	1: $(c', t) \leftarrow c$ 2: $\hat{m} = \text{Dec}(\text{sk}', c')$ 3: $\hat{k} \leftarrow G(\hat{m})$ 4: if $\text{MAC}(\hat{k}, c') = t$ then 5: $K \leftarrow H(\hat{m}, c')$ 6: else 7: $K \leftarrow H(z, c')$ 8: end if 9: return K
$\mathcal{O}^G(m)$	$\mathcal{O}_1^{\text{Decap}}(c)$
1: if $\exists(\tilde{m}, \tilde{k}) \in \mathcal{L}^G : \tilde{m} = m$ then 2: return \tilde{k} 3: end if 4: $k \xleftarrow{\$} \mathcal{K}_{\text{MAC}}$ 5: $\mathcal{L}^G \leftarrow \mathcal{L}^G \cup \{(m, k)\}$ 6: return k	1: $(c', t) \leftarrow c$ 2: if $\exists(\tilde{m}, \tilde{k}) \in \mathcal{L}^G : \tilde{m} = \text{Dec}(\text{sk}', c') \wedge \text{MAC}(\tilde{k}, c') = t$ then 3: $K \leftarrow H(\tilde{m}, c')$ 4: else 5: $K \leftarrow H(z, c')$ 6: end if 7: return K
$\mathcal{O}^H(m, c)$	$\mathcal{O}^H(m, c)$
	1: if $\exists(\tilde{m}, \tilde{c}, \tilde{K}) \in \mathcal{L}^H : \tilde{m} = m \wedge \tilde{c} = c$ then 2: return \tilde{K} 3: end if 4: $K \xleftarrow{\$} \mathcal{K}_{\text{KEM}}$ 5: $\mathcal{L}^H \leftarrow \mathcal{L}^H \cup \{(m, c, K)\}$ 6: return K

Figure 7: Sequence of games

oracle model, such k is a uniformly random sample of \mathcal{K}_{MAC} that the adversary does not know, so for A to produce a valid tag is to produce a forgery against the MAC under an unknown and uniformly random key. Furthermore, the security game does not include a signing oracle, so this is a zero-time forgery. While zero-time forgery is not a standard security definition for a MAC, we can bound it by the advantage of a one-time forgery adversary C :

$$P \left[\mathcal{O}^{\text{Decap}}(c) \neq \mathcal{O}_1^{\text{Decap}}(c) \right] \leq \text{Adv}_{\text{OT-MAC}}(C)$$

Across all q decapsulation queries, the probability that at least one query is a forgery is thus at most $q \cdot P \left[\mathcal{O}^{\text{Decap}}(c) \neq \mathcal{O}_1^{\text{Decap}}(c) \right]$. By the difference lemma:

$$\text{Adv}_{G_0}(A) - \text{Adv}_{G_1}(A) \leq q \cdot \text{Adv}_{\text{OT-MAC}}(C)$$

Game 2 is identical to game 1, except that the challenger samples a uniformly random MAC key $k^* \xleftarrow{\$} \mathcal{K}_{\text{MAC}}$ instead of deriving it from m^* . From A 's perspective the two games are indistinguishable, unless A queries G with the value of m^* . Denote the probability that A queries G with m^* by $P[\text{QUERY } G]$, then:

$$\text{Adv}_{G_1}(A) - \text{Adv}_{G_2}(A) \leq P[\text{QUERY G}]$$

271 *Game 3* is identical to game 2, except that the challenger samples a uniformly random
 272 shared secret $K_0 \xleftarrow{\$} \mathcal{K}_{\text{KEM}}$ instead of deriving it from m^* and c' . From A 's perspective the
 273 two games are indistinguishable, unless A queries H with (m^*, \cdot) . Denote the probability
 274 that A queries H with (m^*, \cdot) by $P[\text{QUERY H}]$, then:

$$\text{Adv}_{G_2}(A) - \text{Adv}_{G_3}(A) \leq P[\text{QUERY H}]$$

275 Since in game 3, both K_0 and K_1 are uniformly random and independent of all other
 276 variables, no adversary can have any advantage: $\text{Adv}_{G_3}(A) = 0$.

$B(\text{pk}, c'^*)$	$\mathcal{O}_B^{\text{Decap}}(c)$
1: $z \xleftarrow{\$} \mathcal{M}$ 2: $k \xleftarrow{\$} \mathcal{K}_{\text{MAC}}$ 3: $t \leftarrow \text{MAC}(k, c'^*)$ 4: $c^* \leftarrow (c'^*, t)$ 5: $K \xleftarrow{\$} \mathcal{K}_{\text{KEM}}$ 6: $\hat{b} \leftarrow A^{\mathcal{O}_B^{\text{Decap}}, \mathcal{O}_B^G, \mathcal{O}_B^H}(\text{pk}, c^*, K)$ 7: if $\text{ABORT}(m)$ then 8: return m 9: end if	1: $(c', t) \leftarrow c$ 2: if $\exists (\tilde{m}, \tilde{k}) \in \mathcal{L}^G : \text{PCO}(c', \tilde{m}) = 1 \wedge \text{MAC}(\tilde{k}, c') = t$ then 3: $K \leftarrow H(\tilde{m}, c')$ 4: else 5: $K \leftarrow H(z, c')$ 6: end if 7: return K
$\mathcal{O}_B^H(m, c)$	$\mathcal{O}_B^G(m)$
if $\text{PCO}(m, c'^*) = 1$ then $\text{ABORT}(m)$ end if if $\exists (\tilde{m}, \tilde{c}, \tilde{K}) \in \mathcal{L}^H : \tilde{m} = m \wedge \tilde{c} = c$ then return \tilde{K} end if $K \xleftarrow{\$} \mathcal{K}_{\text{KEM}}$ $\mathcal{L}^H \leftarrow \mathcal{L}^H \cup \{(m, c, K)\}$ return K	1: if $\text{PCO}(m, c'^*) = 1$ then 2: $\text{ABORT}(m)$ 3: end if 4: if $\exists (\tilde{m}, \tilde{k}) \in \mathcal{L}^G : \tilde{m} = m$ then 5: return \tilde{k} 6: end if 7: $k \xleftarrow{\$} \mathcal{K}_{\text{MAC}}$ 8: $\mathcal{L}^G \leftarrow \mathcal{L}^G \cup \{(m, k)\}$ 9: return k

Figure 8: OW-PCA adversary B simulates game 3 for IND-CCA2 adversary A

277 We will bound $P[\text{QUERY G}]$ and $P[\text{QUERY H}]$ by constructing a OW-PCA adversary B
 278 against the underlying PKE that uses A as a sub-routine. B 's behaviors are summarized
 279 in figure 8.

280 B simulates game 3 for A : receiving the public key pk and challenge encryption c'^* , B
 281 samples random MAC key and session key to produce the challenge encapsulation, then
 282 feeds it to A . When simulating the decapsulation oracle, B uses the plaintext-checking
 283 oracle to look for matching queries in \mathcal{L}^G . When simulating the hash oracles, B uses the
 284 plaintext-checking oracle to detect when $m^* = \text{Dec}(\text{sk}', c'^*)$ has been queried. When m^*
 285 is queried, B terminates A and returns m^* to win the OW-PCA game. In other words:

$$P[\text{QUERY G}] \leq \text{Adv}_{\text{OW-PCA}}(B)$$

$$P[\text{QUERY H}] \leq \text{Adv}_{\text{OW-PCA}}(B)$$

286 Combining all equations above produce the desired security bound. \square

4 Implementation

ML-KEM is an IND-CCA2 secure key encapsulation mechanism standardized by NIST in FIPS 203. The IND-CCA2 security of ML-KEM is achieved in two steps. First, ML-KEM constructs an IND-CPA secure public key encryption scheme K-PKE(KeyGen, Enc, Dec) whose security is based on the conjectured intractability of the module learning with error (MLWE) problems against both classical and quantum adversaries. Then, the U_m^L variant of the Fujisaki-Okamoto transformation [HHK17b] is used to construct the KEM MLKEM(KeyGen, Encap, Decap) by calling K-PKE(KeyGen, Enc, Dec) as sub-routines. Because K-PKE.Enc includes substantially more arithmetics than K-PKE.Dec, by using *re-encryption* and *de-randomization*, ML-KEM’s decapsulation routine incurs significant computational cost.

We implemented the “encrypt-then-MAC” KEM construction using K-PKE as the input PKE and compared its performance against ML-KEM under a variety of scenarios. The experimental data showed that while the “encrypt-then-MAC” construction adds a small amount of computational overhead to the encapsulation routine and a small increase in ciphertext size when compared with ML-KEM, it boasts enormous runtime savings in the decapsulation routine, which makes it particularly suitable for deployment in constrained environment. See appendix 6.1 for comparison with Kyber’s third round submission to NIST’s PQC competition.

We refer readers to [oST24] for the details of the K-PKE routines. The “encrypt-then-MAC” KEM routines are described in figure 9 below.

ML-KEM ⁺ .KeyGen()	ML-KEM ⁺ .Decap(sk, c)
<ol style="list-style-type: none"> 1: $z \xleftarrow{\\$} \{0, 1\}^{256}$ 2: $(pk, sk') \xleftarrow{\\$} \text{K-PKE.KeyGen}()$ 3: $h \leftarrow H(pk)$ 4: $sk \leftarrow (sk' pk h z)$ 5: return (pk, sk) 	<p>Require: Secret key $sk = (sk' pk h z)$</p> <p>Require: Ciphertext $c = (c' t)$</p> <ol style="list-style-type: none"> 1: $(sk', pk, h, z) \leftarrow sk$ 2: $(c', t) \leftarrow c$ 3: $\hat{m} \leftarrow \text{K-PKE.Dec}(sk', c')$ 4: $(\bar{K}, \hat{r}, \hat{k}) \leftarrow \text{XOF}(\hat{m} h)$ 5: $\hat{t} \leftarrow \text{MAC}(\hat{k}, c')$ 6: if $\hat{t} = t$ then 7: $K \leftarrow \text{KDF}(\bar{K} t)$ 8: else 9: $K \leftarrow \text{KDF}(z t)$ 10: end if 11: return K
ML-KEM ⁺ .Encap(pk)	
<p>Require: Public key pk</p> <ol style="list-style-type: none"> 1: $m \xleftarrow{\\$} \{0, 1\}^{256}$ 2: $(\bar{K}, r, k) \leftarrow \text{XOF}(m H(pk))$ 3: $c' \leftarrow \text{K-PKE.Enc}(pk, m, r)$ 4: $t \leftarrow \text{MAC}(k, c')$ 5: $K \leftarrow \text{KDF}(\bar{K} c')$ 6: $c \leftarrow (c', t)$ 7: return (c, K) 	

Figure 9: ML-KEM⁺ routines

Our implementation extended from the reference implementation by the PQCrystals team (<https://github.com/pq-crystals/kyber>). All C code is compiled with GCC 11.4.1 and OpenSSL 3.0.8. All binaries are executed on an AWS c7a.medium instance with an AMD EPYC 9R14 CPU at 3.7 GHz and 1 GB of RAM.

4.1 Choosing a message authenticator

For the ML-KEM⁺ implementation, we instantiated MAC with a selection that covered a wide range of MAC designs, including Poly1305 [Ber05], GMAC [MV04], CMAC [IK03][BR05], and KMAC [Gro13].

Poly1305 and GMAC are both Carter-Wegman style authenticators [WC81] that compute the tag using finite field arithmetic. Generically speaking, Carter-Wegman MAC is parameterized by some finite field \mathbb{F} and the maximal message length $L > 0$. Each symmetric key $k = (k_1, k_2) \xleftarrow{\$} \mathbb{F}^2$ is a pair of uniformly random field elements, and the message is parsed into tuples of field elements up to length L : $m = (m_1, m_2, \dots, m_l) \in \mathbb{F}^{\leq L}$. The tag t is computed by evaluating a polynomial whose coefficients are the message blocks and whose indeterminate is the key:

$$\text{MAC}((k_1, k_2), m) = H_{\text{xpoly}}(k_1, m) + k_2 \quad (1)$$

Where H_{xpoly} is given by:

$$H_{\text{xpoly}}(k_1, m) = k_1^{l+1} + k_1^l \cdot m_1 + k_1^{l-1} \cdot m_2 + \dots + k_1 \cdot m_l$$

The authenticator formulated in equation 1 is a one-time MAC. To make the construction many-time secure, a non-repeating nonce r and a PRF is needed:

$$\text{MAC}((k_1, k_2), m, r) = H_{\text{xpoly}}(k_1, m) \oplus \text{PRF}(k_2, r)$$

Specifically, Poly1305 operates in the prime field \mathbb{F}_q where $q = 2^{130} - 5$ whereas GMAC operates in the binary field $\mathbb{F}_{2^{128}}$. In OpenSSL’s implementation, standalone Poly1305 is a one-time secure MAC, whereas GMAC uses a nonce and AES as the PRF and is thus many-time secure (in OpenSSL GMAC is AES-256-GCM except all data is fed into the “associated data” section and thus not encrypted).

CMAC is based on the CBC-MAC with the block cipher instantiated from AES-256. To compute a CMAC tag, the message is first broke into 128-bit blocks with appropriate padding. Each block is first XOR’d with the previous block’s output, then encrypted under AES using the symmetric key. The final output is XOR’d with a sub key derived from the symmetric key, before being encrypted for one last time.

KMAC is defined in [Gro13] to be based on the family of sponge functions with Keccak permutation as the underlying function. We chose KMAC-256, which uses Shake256 as the underlying extendable output functions. KMAC allows variable-length key and tag, but we chose the 256 bits for key length and 128 bits for tag size for consistency with other authenticators.

To isolate the performance characteristics of each authenticator in our instantiation of ML-KEM⁺, we measured the CPU cycles needed for each authenticator to compute a tag on random inputs whose sizes correspond to the ciphertext sizes of ML-KEM. The measurements are summarized in table 3.

From our testing, we found Poly1305 to exhibit the best performance characteristics. However, there are additional security considerations that may require the use of other less efficient MAC instances. For example, it is possible for an adversary with large computing infrastructure or quantum computers to pre-compute a large lookup table mapping symmetric key to the source plaintext. Upon receiving a ciphertext, the adversary can then search through the lookup table for a matching key, which would’ve revealed the corresponding decryption. We partially mitigated such attack by deriving the symmetric key from both the public key and the plaintext, but in case of a long-term keypair, the adversary might still be able to compute a large lookup table AFTER obtaining the long-term public key. Further mitigation could include using larger-size keys, which can be accomplished either by using a Carter-Wegman MAC that operates on a larger finite field or using a MAC with a variable key-length such as KMAC.

Table 3: CPU cycles needed to compute tag on various input sizes

Input size: 768 bytes			Input size: 1088 bytes			Input size: 1568 bytes		
MAC	Median	Average	MAC	Median	Average	MAC	Median	Average
Poly1305	909	2823	Poly1305	961	2704	Poly1305	1065	1809
GMAC	3899	4859	GMAC	3899	4827	GMAC	4055	5026
CMAC	6291	6373	CMAC	7305	7588	CMAC	8735	8772
KMAC	6373	7791	KMAC	9697	9928	KMAC	11647	12186

4.2 KEM performance

Compared to the U_m^\perp variant of Fujisaki-Okamoto transformed used in ML-KEM, the “encrypt-then-MAC” transformation the following trade-off when given the same input sub-routines:

- Both encapsulation and decapsulation add a small amount of overhead for needing to hash both the PKE plaintext and the PKE ciphertext when deriving the shared secret, where as the U_m^\perp transformation only needs to hash the PKE plaintext.
- The encapsulation routine adds a small amount of run-time overhead for computing the authenticator
- The decapsulation routine enjoys substantial runtime speedup because *re-encryption* is replaced with computing an authenticator
- Ciphertext size increases by the size of an authenticator

Since K-PKE.Enc carries significantly more computational complexity than K-PKE.Dec or any MAC we chose, the performance advantage of the “encrypt-then-MAC” transformation over the U_m^\perp transformation is dominated by the runtime saving gained from replacing *re-encryption* with MAC. A comparison between ML-KEM and variations of the ML-KEM⁺ can be found in table 4

Table 4: CPU cycles of each KEM routine

128-bit security		KEM variant		Encap cycles/tick		Decap cycles/tick	
size parameters (bytes)				Median	Average	Median	Average
pk size	800	ML-KEM-512		91467	92065	121185	121650
sk size	1632	Kyber512		97811	98090	119937	120299
ct size	768	ML-KEM ⁺ -512 w/ Poly1305		93157	93626	33733	33908
KeyGen cycles/tick		ML-KEM ⁺ -512 w/ GMAC		97369	97766	37725	37831
Median	75945	ML-KEM ⁺ -512 w/ CMAC		99739	99959	40117	39943
Average	76171	ML-KEM ⁺ -512 w/ KMAC		101009	101313	40741	40916

192-bit security		KEM variant		Encap cycles/tick		Decap cycles/tick	
size parameters (bytes)				Median	Average	Median	Average
pk size	1184	ML-KEM-768		136405	147400	186445	187529
sk size	2400	Kyber768		153061	153670	182129	182755
ct size	1088	ML-KEM ⁺ -768 w/ Poly1305		146405	146860	43315	43463
KeyGen cycles/tick		ML-KEM ⁺ -768 w/ GMAC		149525	150128	46513	46706
Median	129895	ML-KEM ⁺ -768 w/ CMAC		153139	153735	49841	50074
Average	130650	ML-KEM ⁺ -768 w/ KMAC		155219	155848	52415	52611

256-bit security		KEM variant		Encap cycles/tick		Decap cycles/tick	
size parameters (bytes)				Median	Average	Median	Average
pk size	1568	ML-KEM-1024		199185	199903	246245	247320
sk size	3168	Kyber1024		222351	223260	258231	259067
ct size	1568	ML-KEM ⁺ -1024 w/ Poly1305		205763	206499	51375	51562
KeyGen cycles/tick		ML-KEM ⁺ -1024 w/ GMAC		208805	209681	54573	54780
Median	194921	ML-KEM ⁺ -1024 w/ CMAC		213667	214483	59175	59408
Average	195465	ML-KEM ⁺ -1024 w/ KMAC		216761	217468	62269	62516

4.3 Key exchange protocols

A common application of key encapsulation mechanism is key exchange protocols, where two parties establish a shared secret using a public channel. [BDK⁺18] described three key exchange protocols: unauthenticated key exchange (KE), unilaterally authenticated key exchange (UAKE), and mutually authenticated key exchange (AKE). We instantiated an implementation for each of the three key exchange protocols using different variations of the “encrypt-then-MAC” KEM and compared round trip time with implementations instantiated using ML-KEM.

For clarity, we denote the party who sends the first message to be the client and the other party to be the server. Round trip time (RTT) is defined to be the time interval between the moment before the client starts generating ephemeral keypairs and the moment after the client derives the final session key. All experiments are run on a pair of AWS c7a.medium instances both located in the **us-west-2** region. For each experiment, a total of 10,000 rounds of key exchange are performed, with the median and average round trip time (measured in microsecond) recorded.

4.3.1 Unauthenticated key exchange (KE)

In unauthenticated key exchange, a single pair of ephemeral keypair $(\mathbf{pk}_e, \mathbf{sk}_e) \xleftarrow{\$} \text{KeyGen}()$ is generated by the client. The client transmits the ephemeral public key \mathbf{pk}_e to the server, who runs the encapsulation routine $(c_e, K_e) \xleftarrow{\$} \text{Encap}(\mathbf{pk}_e)$ and transmits the ciphertext c_e back to the client. The client finally decapsulates the ciphertext to recover the shared secret $K_e \leftarrow \text{Decap}(\mathbf{sk}_e, c_e)$. The key exchange routines are summarized in figure 10.

Note that in our implementation, a key derivation function (KDF) is applied to the ephemeral shared secret to derive the final session key. This step is added to maintain consistency with other authenticated key exchange protocols, where the final session key is derived from multiple shared secrets. The key derivation function is instantiated using Shake256, and the final session key is 256 bits in length.

$\text{KE}_c()$	$\text{KE}_s()$
1: $(\mathbf{pk}_e, \mathbf{sk}_e) \xleftarrow{\$} \text{KeyGen}()$	1: $\mathbf{pk}_e \leftarrow \text{read}()$
2: $\text{send}(\mathbf{pk}_e)$	2: $(c_e, K_e) \xleftarrow{\$} \text{Encap}(\mathbf{pk}_e)$
3: $c_e \leftarrow \text{read}()$	3: $\text{send}(c_e)$
4: $K_e \leftarrow \text{Decap}(\mathbf{sk}_e, c_e)$	4: $K \leftarrow \text{KDF}(K_e)$
5: $K \leftarrow \text{KDF}(K)$	5: $\text{return } K$
6: $\text{return } K$	

Figure 10: Unauthenticated key exchange (KE) routines

The RTT comparison is summarized in table 5

Table 5: KE RTT comparison

KEM variant	Client TX bytes	Server TX bytes	RTT time (μs)	
			Median	Average
ML-KEM-512	800	768	92	97
ML-KEM-512 ⁺ w/ Poly1305	800	784	70	72
ML-KEM-512 ⁺ w/ GMAC	800	784	73	76
ML-KEM-512 ⁺ w/ CMAC	800	784	75	79
ML-KEM-512 ⁺ w/ KMAC	800	784	76	78

KEM variant	Client TX bytes	Server TX bytes	RTT time (μs)	
			Median	Average
ML-KEM-768	1184	1088	135	140
ML-KEM-768 ⁺ w/ Poly1305	1184	1104	99	104
ML-KEM-768 ⁺ w/ GMAC	1184	1104	101	105
ML-KEM-768 ⁺ w/ CMAC	1184	1104	103	109
ML-KEM-768 ⁺ w/ KMAC	1184	1104	103	107

KEM variant	Client TX bytes	Server TX bytes	RTT time (μs)	
			Median	Average
ML-KEM-1024	1568	1568	193	199
ML-KEM-1024 ⁺ w/ Poly1305	1568	1584	138	141
ML-KEM-1024 ⁺ w/ GMAC	1568	1584	140	145
ML-KEM-1024 ⁺ w/ CMAC	1568	1584	143	148
ML-KEM-1024 ⁺ w/ KMAC	1568	1584	144	149

4.3.2 Unilaterally authenticated key exchange (UAKE)

In unilaterally authenticated key exchange, the authenticating party proves its identity to the other party by demonstrating possession of a secret key that corresponds to a published long-term public key. In this implementation, the client possesses the long-term public key \mathbf{pk}_S of the server, and the server authenticates itself by demonstrating possession of the corresponding long-term secret key \mathbf{sk}_S . UAKE routines are summarized in figure 11.

In addition to the long-term key, the client will also generate an ephemeral keypair as it does in an unauthenticated key exchange, and the session key is derived by applying the KDF to the concatenation of both the ephemeral shared secret and the shared secret encapsulated under server’s long-term key. This helps the key exchange to achieve weak forward secrecy (citation needed).

Using KEM for authentication is especially interesting within the context of post-quantum cryptography: post-quantum KEM schemes usually enjoy better performance characteristics than post-quantum signature schemes with faster runtime, smaller memory footprint, and smaller communication sizes. KEMTLS was proposed in 2020 as an alternative to existing TLS handshake protocols, and many experimental implementations have demonstrated the performance advantage. (citation needed).

$\text{UAKE}_c(\text{pk}_S)$	$\text{UAKE}_s(\text{sk}_S)$
Require: Server's long-term pk_S	Require: Server's long-term sk_S
1: $(\text{pk}_e, \text{sk}_e) \xleftarrow{\$} \text{KeyGen}()$	1: $(\text{pk}_e, c_S) \leftarrow \text{read}()$
2: $(c_S, K_S) \xleftarrow{\$} \text{Encap}(\text{pk}_S)$	2: $K_S \leftarrow \text{Decap}(\text{sk}_S, c_S)$
3: $\text{send}(\text{pk}_e, c_S)$	3: $(c_e, K_e) \xleftarrow{\$} \text{Encap}(\text{pk}_e)$
4: $c_e \leftarrow \text{read}()$	4: $\text{send}(c_e)$
5: $K_e \leftarrow \text{Decap}(\text{sk}_e, c_e)$	5: $K \leftarrow \text{KDF}(K_e \ K_S)$
6: $K \leftarrow \text{KDF}(K_e \ K_S)$	6: return K
7: return K	

Figure 11: Unilaterally authenticated key exchange (UAKE) routines

Table 6: UAKE RTT comparison

KEM variant	Client TX bytes	Server TX bytes	RTT time (μs)	
			Median	Average
ML-KEM-512	1568	768	145	151
ML-KEM-512 ⁺ w/ Poly1305	1584	784	103	106
ML-KEM-512 ⁺ w/ GMAC	1584	784	106	110
ML-KEM-512 ⁺ w/ CMAC	1584	784	108	112
ML-KEM-512 ⁺ w/ KMAC	1584	784	109	113

KEM variant	Client TX bytes	Server TX bytes	RTT time (μs)	
			Median	Average
ML-KEM-768	2272	1088	215	222
ML-KEM-768 ⁺ w/ Poly1305	2288	1104	144	150
ML-KEM-768 ⁺ w/ GMAC	2288	1104	149	156
ML-KEM-768 ⁺ w/ CMAC	2288	1104	153	160
ML-KEM-768 ⁺ w/ KMAC	2288	1104	154	159

KEM variant	Client TX bytes	Server TX bytes	RTT time (μs)	
			Median	Average
ML-KEM-1024	3136	1568	310	318
ML-KEM-1024 ⁺ w/ Poly1305	3152	1584	202	209
ML-KEM-1024 ⁺ w/ GMAC	3152	1584	212	228
ML-KEM-1024 ⁺ w/ CMAC	3152	1584	212	218
ML-KEM-1024 ⁺ w/ KMAC	3152	1584	213	220

4.3.3 Mutually authenticated key exchange (AKE)

Mutually authenticated key exchange is largely identical to unilaterally authenticated key exchange, except for that client authentication is required. This means that client possesses server's long-term public key and its own long-term secret key, while the server possesses client's long-term public key and its own long-term secret key. The session key is derived by applying KDF onto the concatenation of shared secrets produced under the ephemeral keypair, server's long-term keypair, and client's long-term keypair, in this order.

AKE_C(pk_S, sk_C)	AKE_S(sk_S, pk_C)
Require: Server's long-term pk _S	Require: Server's long-term sk _S
Require: Client's long-term sk _C	Require: Client's long-term pk _C
1: (pk _e , sk _e) $\xleftarrow{\$}$ KeyGen()	1: (pk _e , c _S) \leftarrow read()
2: (c _S , K _S) $\xleftarrow{\$}$ Encap(pk _S)	2: K _S \leftarrow Decap(sk _S , c _S)
3: send(pk _e , c _S)	3: (c _e , K _e) $\xleftarrow{\$}$ Encap(pk _e)
4: (c _e , c _C) \leftarrow read()	4: (c _C , K _C) $\xleftarrow{\$}$ Encap(pk _C)
5: K _e \leftarrow Decap(sk _e , c _e)	5: send(c _e , c _C)
6: K _C \leftarrow Decap(sk _e , c _C)	6: K \leftarrow KDF(K _e K _S K _C)
7: K \leftarrow KDF(K _e K _S K _C)	7: return K
8: return K	

Figure 12: Mutually authenticated key exchange (AKE) routines

Table 7: AKE RTT comparison

KEM variant	Client TX bytes	Server TX bytes	RTT time (μ s)	
			Median	Average
ML-KEM-512	1568	1536	220	213
ML-KEM-512 ⁺ w/ Poly1305	1584	1568	133	138
ML-KEM-512 ⁺ w/ GMAC	1584	1568	139	143
ML-KEM-512 ⁺ w/ CMAC	1584	1568	143	148
ML-KEM-512 ⁺ w/ KMAC	1584	1568	145	151

KEM variant	Client TX bytes	Server TX bytes	RTT time (μ s)	
			Median	Average
ML-KEM-768	2272	2176	294	301
ML-KEM-768 ⁺ w/ Poly1305	2288	2208	190	196
ML-KEM-768 ⁺ w/ GMAC	2288	2208	197	210
ML-KEM-768 ⁺ w/ CMAC	2288	2208	202	208
ML-KEM-768 ⁺ w/ KMAC	2288	2208	204	210

KEM variant	Client TX bytes	Server TX bytes	RTT time (μ s)	
			Median	Average
ML-KEM-1024	3136	3136	512	511
ML-KEM-1024 ⁺ w/ Poly1305	3152	3168	266	273
ML-KEM-1024 ⁺ w/ GMAC	3152	3168	273	282
ML-KEM-1024 ⁺ w/ CMAC	3152	3168	280	287
ML-KEM-1024 ⁺ w/ KMAC	3152	3168	282	288

5 Conclusions and future works

The “encrypt-then-MAC” transformation is a generic KEM construction that achieves IND-CCA2 security under the random oracle model if the input PKE is OW-PCA secure. Compared to the Fujisaki-Okamoto transformation, our construction replaced *de-randomization* and *re-encryption* with a message authenticator. At the cost of some minimal increase in communication size and encapsulation runtime, our construction achieves significant efficiency gains in the decapsulation routine. In practical key exchange protocols, our construction saves between 35-45% in round trip time.

Other suitable post-quantum instantiation. In section 4, we instantiated the “encrypt-then-MAC” KEM transformation with the PKE subroutines of ML-KEM. Unfortunately, ML-KEM is vulnerable to plaintext checking attacks that can recover complete

secret key if the keypair is reused [BDH⁺19]. In fact, as Chris Peikert pointed out, because of the search-decision equivalence of LWE problems, almost all LWE-based cryptosystems will be vulnerable to plaintext-checking attacks [Pei14]. This naturally raises the question of finding other post-quantum primitives to instantiate “encrypt-then-MAC” with, such as other lattice-based cryptosystems not based on LWE, code-based cryptosystems and isogeny-based cryptosystems.

References

- [ABC⁺20] Martin R. Albrecht, Daniel J. Bernstein, Tung Chou, Carlos Cid, Jan Gilcher, Tanja Lange, Varun Maram, Ingo von Maurich, Rafael Misoczki, Ruben Niederhagen, Kenneth G. Paterson, Edoardo Persichetti, Christiane Peters, Peter Schwabe, Nicolas Sendrier, Jakub Szefer, Cen Jung Tjhai, Martin Tomlinson, and Wen Wang. Classic mceliece. Technical report, National Institute of Standards and Technology, 2020.
- [ABD⁺19] Roberto Avanzi, Joppe Bos, Léo Ducas, Eike Kiltz, Tancrede Lepoint, Vadim Lyubashevsky, John M Schanck, Peter Schwabe, Gregor Seiler, and Damien Stehlé. Crystals-kyber algorithm specifications and supporting documentation. *NIST PQC Round*, 2(4):1–43, 2019.
- [ABR99] Michel Abdalla, Mihir Bellare, and Phillip Rogaway. DHAES: an encryption scheme based on the diffie-hellman problem. *IACR Cryptol. ePrint Arch.*, page 7, 1999.
- [ABR01] Michel Abdalla, Mihir Bellare, and Phillip Rogaway. The oracle diffie-hellman assumptions and an analysis of DHIES. In David Naccache, editor, *Topics in Cryptology - CT-RSA 2001, The Cryptographer’s Track at RSA Conference 2001, San Francisco, CA, USA, April 8-12, 2001, Proceedings*, volume 2020 of *Lecture Notes in Computer Science*, pages 143–158. Springer, 2001.
- [BCD⁺16] Joppe W. Bos, Craig Costello, Léo Ducas, Ilya Mironov, Michael Naehrig, Valeria Nikolaenko, Ananth Raghunathan, and Douglas Stebila. Frodo: Take off the ring! practical, quantum-secure key exchange from LWE. In Edgar R. Weippl, Stefan Katzenbeisser, Christopher Kruegel, Andrew C. Myers, and Shai Halevi, editors, *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security, Vienna, Austria, October 24-28, 2016*, pages 1006–1018. ACM, 2016.
- [BDH⁺19] Ciprian Baetu, F. Betül Durak, Loïs Huguenin-Dumittan, Abdullah Talayhan, and Serge Vaudenay. Misuse attacks on post-quantum cryptosystems. In Yuval Ishai and Vincent Rijmen, editors, *Advances in Cryptology - EUROCRYPT 2019 - 38th Annual International Conference on the Theory and Applications of Cryptographic Techniques, Darmstadt, Germany, May 19-23, 2019, Proceedings, Part II*, volume 11477 of *Lecture Notes in Computer Science*, pages 747–776. Springer, 2019.
- [BDK⁺18] Joppe W. Bos, Léo Ducas, Eike Kiltz, Tancrede Lepoint, Vadim Lyubashevsky, John M. Schanck, Peter Schwabe, Gregor Seiler, and Damien Stehlé. CRYSTALS - kyber: A cca-secure module-lattice-based KEM. In *2018 IEEE European Symposium on Security and Privacy, EuroS&P 2018, London, United Kingdom, April 24-26, 2018*, pages 353–367. IEEE, 2018.
- [BDPR98] Mihir Bellare, Anand Desai, David Pointcheval, and Phillip Rogaway. Relations among notions of security for public-key encryption schemes. In Hugo Krawczyk,

- 483 editor, *Advances in Cryptology - CRYPTO '98, 18th Annual International*
484 *Cryptology Conference, Santa Barbara, California, USA, August 23-27, 1998,*
485 *Proceedings*, volume 1462 of *Lecture Notes in Computer Science*, pages 26–45.
486 Springer, 1998.
- 487 [Ber05] Daniel J. Bernstein. The poly1305-aes message-authentication code. In Henri
488 Gilbert and Helena Handschuh, editors, *Fast Software Encryption: 12th Inter-*
489 *national Workshop, FSE 2005, Paris, France, February 21-23, 2005, Revised*
490 *Selected Papers*, volume 3557 of *Lecture Notes in Computer Science*, pages
491 32–49. Springer, 2005.
- 492 [Ble98] Daniel Bleichenbacher. Chosen ciphertext attacks against protocols based on
493 the RSA encryption standard PKCS #1. In Hugo Krawczyk, editor, *Advances*
494 *in Cryptology - CRYPTO '98, 18th Annual International Cryptology Conference,*
495 *Santa Barbara, California, USA, August 23-27, 1998, Proceedings*, volume 1462
496 of *Lecture Notes in Computer Science*, pages 1–12. Springer, 1998.
- 497 [BN00] Mihir Bellare and Chanathip Namprempre. Authenticated encryption: Re-
498 lations among notions and analysis of the generic composition paradigm. In
499 Tatsuaki Okamoto, editor, *Advances in Cryptology - ASIACRYPT 2000, 6th*
500 *International Conference on the Theory and Application of Cryptology and*
501 *Information Security, Kyoto, Japan, December 3-7, 2000, Proceedings*, volume
502 1976 of *Lecture Notes in Computer Science*, pages 531–545. Springer, 2000.
- 503 [BP18] Daniel J. Bernstein and Edoardo Persichetti. Towards KEM unification. *IACR*
504 *Cryptol. ePrint Arch.*, page 526, 2018.
- 505 [BR94] Mihir Bellare and Phillip Rogaway. Optimal asymmetric encryption. In Al-
506 fredo De Santis, editor, *Advances in Cryptology - EUROCRYPT '94, Workshop*
507 *on the Theory and Application of Cryptographic Techniques, Perugia, Italy,*
508 *May 9-12, 1994, Proceedings*, volume 950 of *Lecture Notes in Computer Science*,
509 pages 92–111. Springer, 1994.
- 510 [BR05] John Black and Phillip Rogaway. CBC macs for arbitrary-length messages:
511 The three-key constructions. *J. Cryptol.*, 18(2):111–131, 2005.
- 512 [CHJ⁺02] Jean-Sébastien Coron, Helena Handschuh, Marc Joye, Pascal Paillier, David
513 Pointcheval, and Christophe Tymen. GEM: A generic chosen-ciphertext secure
514 encryption method. In Bart Preneel, editor, *Topics in Cryptology - CT-RSA*
515 *2002, The Cryptographer's Track at the RSA Conference, 2002, San Jose, CA,*
516 *USA, February 18-22, 2002, Proceedings*, volume 2271 of *Lecture Notes in*
517 *Computer Science*, pages 263–276. Springer, 2002.
- 518 [Den03] Alexander W. Dent. A designer's guide to kems. In Kenneth G. Paterson, editor,
519 *Cryptography and Coding, 9th IMA International Conference, Cirencester, UK,*
520 *December 16-18, 2003, Proceedings*, volume 2898 of *Lecture Notes in Computer*
521 *Science*, pages 133–151. Springer, 2003.
- 522 [DKRV18] Jan-Pieter D'Anvers, Angshuman Karmakar, Sujoy Sinha Roy, and Frederik
523 Vercauteren. Saber: Module-lwr based key exchange, cpa-secure encryption
524 and cca-secure KEM. In Antoine Joux, Abderrahmane Nitaj, and Tajjeeddine
525 Rachidi, editors, *Progress in Cryptology - AFRICACRYPT 2018 - 10th Inter-*
526 *national Conference on Cryptology in Africa, Marrakesh, Morocco, May 7-9,*
527 *2018, Proceedings*, volume 10831 of *Lecture Notes in Computer Science*, pages
528 282–305. Springer, 2018.

- [DNR04] Cynthia Dwork, Moni Naor, and Omer Reingold. Immunizing encryption schemes from decryption errors. In Christian Cachin and Jan Camenisch, editors, *Advances in Cryptology - EUROCRYPT 2004, International Conference on the Theory and Applications of Cryptographic Techniques, Interlaken, Switzerland, May 2-6, 2004, Proceedings*, volume 3027 of *Lecture Notes in Computer Science*, pages 342–360. Springer, 2004.
- [FO99] Eiichiro Fujisaki and Tatsuaki Okamoto. Secure integration of asymmetric and symmetric encryption schemes. In Michael J. Wiener, editor, *Advances in Cryptology - CRYPTO '99, 19th Annual International Cryptology Conference, Santa Barbara, California, USA, August 15-19, 1999, Proceedings*, volume 1666 of *Lecture Notes in Computer Science*, pages 537–554. Springer, 1999.
- [FO13] Eiichiro Fujisaki and Tatsuaki Okamoto. Secure integration of asymmetric and symmetric encryption schemes. *J. Cryptol.*, 26(1):80–101, 2013.
- [FOPS01] Eiichiro Fujisaki, Tatsuaki Okamoto, David Pointcheval, and Jacques Stern. RSA-OAEP is secure under the RSA assumption. In Joe Kilian, editor, *Advances in Cryptology - CRYPTO 2001, 21st Annual International Cryptology Conference, Santa Barbara, California, USA, August 19-23, 2001, Proceedings*, volume 2139 of *Lecture Notes in Computer Science*, pages 260–274. Springer, 2001.
- [Gro13] Joint Task Force Transformation Initiative Interagency Working Group. Security and privacy controls for federal information systems and organizations. Technical Report NIST Special Publication (SP) 800-53, Rev. 4, Includes updates as of January 22, 2015, National Institute of Standards and Technology, Gaithersburg, MD, 2013.
- [HHK17a] Dennis Hofheinz, Kathrin Hövelmanns, and Eike Kiltz. A modular analysis of the fujisaki-okamoto transformation. In Yael Kalai and Leonid Reyzin, editors, *Theory of Cryptography - 15th International Conference, TCC 2017, Baltimore, MD, USA, November 12-15, 2017, Proceedings, Part I*, volume 10677 of *Lecture Notes in Computer Science*, pages 341–371. Springer, 2017.
- [HHK17b] Dennis Hofheinz, Kathrin Hövelmanns, and Eike Kiltz. A modular analysis of the fujisaki-okamoto transformation. In Yael Kalai and Leonid Reyzin, editors, *Theory of Cryptography - 15th International Conference, TCC 2017, Baltimore, MD, USA, November 12-15, 2017, Proceedings, Part I*, volume 10677 of *Lecture Notes in Computer Science*, pages 341–371. Springer, 2017.
- [HHM22] Kathrin Hövelmanns, Andreas Hülsing, and Christian Majenz. Failing gracefully: Decryption failures and the fujisaki-okamoto transform. In Shweta Agrawal and Dongdai Lin, editors, *Advances in Cryptology - ASIACRYPT 2022 - 28th International Conference on the Theory and Application of Cryptology and Information Security, Taipei, Taiwan, December 5-9, 2022, Proceedings, Part IV*, volume 13794 of *Lecture Notes in Computer Science*, pages 414–443. Springer, 2022.
- [IK03] Tetsu Iwata and Kaoru Kurosawa. OMAC: one-key CBC MAC. In Thomas Johansson, editor, *Fast Software Encryption, 10th International Workshop, FSE 2003, Lund, Sweden, February 24-26, 2003, Revised Papers*, volume 2887 of *Lecture Notes in Computer Science*, pages 129–153. Springer, 2003.
- [Kal98] Burt Kaliski. PKCS #1: RSA encryption version 1.5. *RFC*, 2313:1–19, 1998.

- [Kra01] Hugo Krawczyk. The order of encryption and authentication for protecting communications (or: How secure is ssl?). In Joe Kilian, editor, *Advances in Cryptology - CRYPTO 2001, 21st Annual International Cryptology Conference, Santa Barbara, California, USA, August 19-23, 2001, Proceedings*, volume 2139 of *Lecture Notes in Computer Science*, pages 310–331. Springer, 2001.
- [MKJR16] Kathleen Moriarty, Burt Kaliski, Jakob Jonsson, and Andreas Rusch. PKCS #1: RSA Cryptography Specifications Version 2.2. RFC 8017, November 2016.
- [MV04] David A. McGrew and John Viega. The security and performance of the galois/counter mode (GCM) of operation. In Anne Canteaut and Kapalee Viswanathan, editors, *Progress in Cryptology - INDOCRYPT 2004, 5th International Conference on Cryptology in India, Chennai, India, December 20-22, 2004, Proceedings*, volume 3348 of *Lecture Notes in Computer Science*, pages 343–355. Springer, 2004.
- [OP01a] Tatsuaki Okamoto and David Pointcheval. The gap-problems: A new class of problems for the security of cryptographic schemes. In Kwangjo Kim, editor, *Public Key Cryptography, 4th International Workshop on Practice and Theory in Public Key Cryptography, PKC 2001, Cheju Island, Korea, February 13-15, 2001, Proceedings*, volume 1992 of *Lecture Notes in Computer Science*, pages 104–118. Springer, 2001.
- [OP01b] Tatsuaki Okamoto and David Pointcheval. REACT: rapid enhanced-security asymmetric cryptosystem transform. In David Naccache, editor, *Topics in Cryptology - CT-RSA 2001, The Cryptographer’s Track at RSA Conference 2001, San Francisco, CA, USA, April 8-12, 2001, Proceedings*, volume 2020 of *Lecture Notes in Computer Science*, pages 159–175. Springer, 2001.
- [oST24] National Institute of Standards and Technology. Module-lattice-based key-encapsulation mechanism standard. Technical Report Federal Information Processing Standards Publication (FIPS) NIST FIPS 203, U.S. Department of Commerce, Washington, D.C., 2024.
- [Pei14] Chris Peikert. Lattice cryptography for the internet. In Michele Mosca, editor, *Post-Quantum Cryptography - 6th International Workshop, PQCrypto 2014, Waterloo, ON, Canada, October 1-3, 2014, Proceedings*, volume 8772 of *Lecture Notes in Computer Science*, pages 197–219. Springer, 2014.
- [RRCB19] Prasanna Ravi, Sujoy Sinha Roy, Anupam Chattopadhyay, and Shivam Bhasin. Generic side-channel attacks on cca-secure lattice-based PKE and KEM schemes. *IACR Cryptol. ePrint Arch.*, page 948, 2019.
- [RSA78] Ronald L. Rivest, Adi Shamir, and Leonard M. Adleman. A method for obtaining digital signatures and public-key cryptosystems. *Commun. ACM*, 21(2):120–126, 1978.
- [Sho01] Victor Shoup. A proposal for an ISO standard for public key encryption. *IACR Cryptol. ePrint Arch.*, page 112, 2001.
- [Sho02] Victor Shoup. OAEP reconsidered. *J. Cryptol.*, 15(4):223–249, 2002.
- [Sho04] Victor Shoup. Sequences of games: a tool for taming complexity in security proofs. *IACR Cryptol. ePrint Arch.*, page 332, 2004.

- [TUX⁺23] Yutaro Tanaka, Rei Ueno, Keita Xagawa, Akira Ito, Junko Takahashi, and Naofumi Homma. Multiple-valued plaintext-checking side-channel attacks on post-quantum kems. *IACR Trans. Cryptogr. Hardw. Embed. Syst.*, 2023(3):473–503, 2023.
- [UXT⁺22] Rei Ueno, Keita Xagawa, Yutaro Tanaka, Akira Ito, Junko Takahashi, and Naofumi Homma. Curse of re-encryption: A generic power/em analysis on post-quantum kems. *IACR Trans. Cryptogr. Hardw. Embed. Syst.*, 2022(1):296–322, 2022.
- [WC81] Mark N. Wegman and Larry Carter. New hash functions and their use in authentication and set equality. *J. Comput. Syst. Sci.*, 22(3):265–279, 1981.

6 Appendix

6.1 Performance comparison between ML-KEM⁺ and Kyber

ML-KEM directly evolved from CRYSTALS-Kyber’s third round submission to NIST’s post quantum cryptography competition. While their IND-CPA subroutines (see figure 13) are identical, ML-KEM deviated from Kyber by choosing a different variant of the Fujisaki-Okamoto transformation.

K-PKE.KeyGen()	K-PKE.Enc(pk, m)	K-PKE.Dec(sk, c)
1: $A \xleftarrow{\$} R_q^{k \times k}$	Ensure: $\text{pk} = (A, \mathbf{t})$	Ensure: $c = (c_1, c_2)$
2: $\mathbf{s} \xleftarrow{\$} \mathcal{X}_{\eta_1}^k$	Ensure: $m \in R_2$	Ensure: $\text{sk} = \mathbf{s}$
3: $\mathbf{e} \xleftarrow{\$} \mathcal{X}_{\eta_1}^k$	1: $\mathbf{r} \xleftarrow{\$} \mathcal{X}_{\eta_1}^k$	1: $\hat{m} \leftarrow c_2 - \mathbf{c}_1^\top \cdot \mathbf{s}$
4: $\mathbf{t} \leftarrow A\mathbf{s} + \mathbf{e}$	2: $\mathbf{e}_1 \xleftarrow{\$} \mathcal{X}_{\eta_2}^k$	2: $\hat{m} \leftarrow \text{Round}(\hat{m})$
5: $\text{pk} \leftarrow (A, \mathbf{t})$	3: $e_2 \xleftarrow{\$} \mathcal{X}_{\eta_2}$	3: return \hat{m}
6: $\text{sk} \leftarrow \mathbf{s}$	4: $\mathbf{c}_1 \leftarrow A\mathbf{r} + \mathbf{e}_1$	
7: return (pk, sk)	5: $c_2 \leftarrow \mathbf{t}^\top \mathbf{r} + e_2 + m \cdot \lfloor \frac{q}{2} \rfloor$	
	6: return (c ₁ , c ₂)	

Figure 13: K-PKE routines are identical between Kyber and ML-KEM

CRYSTALS-Kyber uses the U^\perp variant, where the shared secret is derived from both the plaintext and the ciphertext. On the other hand, because by using *re-encryption* and *de-randomization*, the PKE is already made *rigid*, the CRYSTALS-Kyber team decided to use the U_m^\perp variant, where the shared secret is derived from the plaintext alone.

KEM.KeyGen()	KEM.Decap(sk, c)
1: $z \xleftarrow{\$} \{0, 1\}^{256}$ 2: $(pk, sk') \xleftarrow{\$} \text{PKE.KeyGen}()$ 3: $sk \leftarrow (sk' pk H(pk) z)$ 4: return (pk, sk)	Ensure: $sk = (sk' pk H(pk) z)$ 1: $\hat{m} \leftarrow \text{PKE.Dec}(sk', c)$ 2: $(\bar{K}, \hat{r}) \leftarrow G(\hat{m} H(pk))$ 3: if $\text{PKE.Enc}(pk, \hat{m}, \hat{r}) = c$ then 4: $K \leftarrow \text{KDF}(\bar{K}, H(c)) \quad \triangleright U^\mathcal{X}$ 5: $K \leftarrow \bar{K} \quad \triangleright U_m^\mathcal{X}$ 6: else 7: $K \leftarrow \text{KDF}(z H(c))$ 8: end if 9: return K
KEM.Encap(pk)	
1: $m \xleftarrow{\$} \{0, 1\}^{256}$ 2: $(\bar{K}, r) \leftarrow G(m H(pk))$ 3: $c \leftarrow \text{PKE.Enc}(pk, m, r)$ 4: $K \leftarrow \text{KDF}(\bar{K} H(c)) \quad \triangleright U^\mathcal{X}$ 5: $K \leftarrow \bar{K} \quad \triangleright U_m^\mathcal{X}$ 6: return (c, K)	

Figure 14: Kyber uses $U^\mathcal{X}$ variant. ML-KEM uses $U_m^\mathcal{X}$ variant.

638 The reason for ML-KEM to use a different variant of the Fujisaki-Okamoto trans-
639 formation is two-fold. The first reason is performance: using the $U_m^\mathcal{X}$ trans-
640 formation saves the need to hash the ciphertext, and since Kyber/ML-KEM's performance is mainly
641 bottlenecked by the symmetric components, omitting the hash leads to significant runtime
642 savings (up to 17% in AVX-2 optimized implementations). The second reason is the
643 simplified security proof and tighter security bounds of the $U_m^\mathcal{X}$ variant compared to the
644 $U^\mathcal{X}$ variant. We will omit the details of the security proof and refer readers to [HHK17b].

645 In section 4, we mainly compared ML-KEM⁺ with ML-KEM, but we would like to point
646 out that, because Kyber uses the $U^\mathcal{X}$ variant and needs to hash the ciphertext for deriving
647 the shared secret, the performance advantage of ML-KEM⁺ over Kyber will be even greater.