

Fast Fujisaki-Okamoto transformation using encrypt-then-mac and applications to Kyber

Anonymous Submission

Abstract. The modular Fujisaki-Okamoto (FO) transformation takes public-key encryption with weaker security and constructs a key encapsulation mechanism (KEM) with indistinguishability under adaptive chosen ciphertext attacks. While the modular FO transform enjoys tight security bound and quantum resistance, it also suffers from computational inefficiency due to using de-randomization and re-encryption for providing ciphertext integrity. In this work, we propose an alternative modular FO transformation that replaces re-encryption with a message authentication code (MAC) and prove the security bound of our construction. We then instantiate a concrete instance with ML-KEM and show that when re-encryption incurs significant computational cost, our construction provides substantial runtime speedup and reduced memory footprint.

Keywords: Key encapsulation mechanism, post-quantum cryptography, lattice cryptography, Fujisaki-Okamoto transformation

1 Introduction

The Fujisaki-Okamoto transformation [FO99] is a generic construction that takes cryptographic primitives of lesser security and constructs a public-key encryption scheme with indistinguishability under adaptive chosen ciphertext attacks. Later works extended the original transformation to the construction of key encapsulation mechanism, which has been adopted by many post-quantum schemes such as Kyber [BDK⁺18] (standardized by NIST into ML-KEM [KE23]).

The current state of the FO transformation enjoys tight security bound and quantum resistance [HHK17], but also leaves many open questions. One such problem is the use of re-encryption for providing ciphertext integrity [BP18], which requires the decryption/decapsulation to run the encryption routine as a subroutine. In many post-quantum schemes, such as Kyber, the encryption routine is substantially computationally more expensive than the decryption routine.

The problem of ciphertext integrity was solved in symmetric cryptography. Given a semantically secure symmetric cipher and an existentially unforgeable message authentication code, combining them using “encrypt-then-mac” provides authenticated encryption [BN00]. We took inspiration from this strategy and applied a similar technique to provide ciphertext integrity for a public-key encryption scheme, which then translates to an IND-CCA secure KEM. Using a message authentication code for ciphertext integrity replaces the re-encryption step in decryption with the computation of a tag, which should offer significant performance improvements while maintaining comparable level of security.

The main challenge in applying “encrypt-then-mac” to public-key cryptography is the lack of a pre-shared MAC key. We proposed to derive the shared MAC key by hashing the plaintext message. We will prove in section 3 that, under the random oracle model, the MAC key is securely hidden behind the hash function, and producing a valid pair of ciphertext and tag without full knowledge of the plaintext constitutes a forgery attack on the message authentication code. Thanks to the modular construction in [HHK17],

providing ciphertext integrity in the underlying encryption scheme gives us an IND-CCA secure KEM for free.

In section 4.1, we instantiate concrete instances of our proposed transformation by modifying ML-KEM. We will demonstrate that, at the cost of small increase in encryption runtime and ciphertext size, our construction reduces both the runtime and memory footprint of the decryption routine.

2 Preliminaries and previous results

2.1 Public-key encryption scheme

We define a public key encryption scheme PKE to be a collection of three routines ($\text{Gen}, \text{Enc}, \text{Dec}$) defined over a finite message space \mathcal{M} and some ciphertext space \mathcal{C} . Many encryption routines are probabilistic, and we define their source of randomness to come from some coin space \mathcal{R} .

The encryption routine $\text{Enc}(\text{pk}, m)$ takes a public key, a plaintext message, and outputs a ciphertext $c \in \mathcal{C}$. Where the encryption routine is probabilistic, specifying a pseudorandom seed $r \in \mathcal{R}$ will make the encryption routine behave deterministically. The decryption routine $\text{Dec}(\text{sk}, c)$ takes a secret key, a ciphertext, and outputs the decryption \hat{m} if the ciphertext is valid under the given secret key, or the rejection symbol \perp if the ciphertext is invalid.

2.1.1 Correctness

It is common to require a PKE to be perfectly correct, meaning that for all possible keypairs (pk, sk) and plaintext messages $m \in \mathcal{M}$, $\text{Dec}(\text{sk}, \text{Enc}(\text{pk}, m)) = m$ at all times. However, some encryption schemes, including many popular lattice-based schemes, admit a non-zero probability of decryption failure: $\text{Dec}(\text{sk}, \text{Enc}(\text{pk}, m)) \neq m$. Furthermore, [HHK17] and [ABD⁺19] explained how decryption failure played a role in an adversary's advantage. In this paper, we inherit the definition for correctness from [HHK17]:

Definition 1 (δ -correctness). A public key encryption scheme PKE is δ -correct if

$$\mathbf{E}[\max_{m \in \mathcal{M}} P[\text{Dec}(\text{sk}, c) \neq m \mid c \xleftarrow{\$} \text{Enc}(\text{pk}, m)]] \leq \delta$$

Where the expectation is taken over the probability distribution of keypairs $(\text{pk}, \text{sk}) \xleftarrow{\$} \text{Gen}()$

2.1.2 Security

We discuss the security of a PKE using the sequence of games described in [Sho04]. Specifically, we first define the OW-ATK and the IND-CPA game as they pertain to a public key encryption scheme. In later section we will define the IND-CCA game as it pertains to a key encapsulation mechanism.

In the OW-ATK game, an adversary's goal is to recover the decryption of a randomly generated ciphertext.

The adversary \mathcal{A} with access to oracle(s) \mathcal{O}_{ATK} wins the game if its guess \hat{m} is equal to the challenge plaintext m^* . The *advantage* $\epsilon_{\text{OW-ATK}}$ of an adversary in this game is the probability that it wins the game.

The choice of oracle(s) \mathcal{O}_{ATK} depends on the choice of ATK. Specifically:

Algorithm 1 The OW-ATK game

```

1:  $(\mathbf{pk}, \mathbf{sk}) \xleftarrow{\$} \text{Gen}(1^\lambda)$ 
2:  $m^* \xleftarrow{\$} \mathcal{M}$ 
3:  $c^* \xleftarrow{\$} \text{Enc}(\mathbf{pk}, m^*)$ 
4:  $\hat{m} \xleftarrow{\$} \mathcal{A}^{\mathcal{O}_{\text{ATK}}}(1^\lambda, \mathbf{pk}, c^*)$ 
5: return  $\llbracket m^* = \hat{m} \rrbracket$ 

```

Figure 1: The OW-ATK game**Algorithm 2** PCO($m \in \mathcal{M}, c \in \mathcal{C}$)

```

1: return  $\llbracket \text{Dec}(\mathbf{sk}, c) = m \rrbracket$ 

```

Algorithm 3 CVO($c \in \mathcal{C}$)

```

1: return  $\llbracket \text{Dec}(\mathbf{sk}, c) \in \mathcal{M} \rrbracket$ 

```

Figure 2: The Plaintext-Checking Oracle PCO**Figure 3:** the Ciphertext-Validation Oracle CVO

$$\mathcal{O}_{\text{ATK}} = \begin{cases} - & \text{ATK} = \text{CPA} \\ \text{PCO} & \text{ATK} = \text{PCA} \\ \text{CVO} & \text{ATK} = \text{VA} \\ \text{PCO}, \text{CVO} & \text{ATK} = \text{PCVA} \end{cases}$$

83 Where the definitions of plaintext-checking oracle PCO and the ciphertext-validation
84 oracle CVO are inherited from [HHK17]

85 In the IND-CPA game (algorithm 4), an adversary's goal is to distinguish the encryption
86 of one message from the encryption of another message. Given the public key, the adversary
87 outputs two adversarially chosen messages and obtains the encryption of a random choice
88 between these two messages. The adversary wins the IND-CPA game if it correctly identifies
89 which message the encryption is obtained from.

Algorithm 4 The IND-CPA game

```

1:  $(\mathbf{pk}, \mathbf{sk}) \xleftarrow{\$} \text{Gen}(1^\lambda)$ 
2:  $(m_0, m_1) \xleftarrow{\$} \mathcal{A}(1^\lambda, \mathbf{pk})$ 
3:  $b \xleftarrow{\$} \{0, 1\}$ 
4:  $c^* \xleftarrow{\$} \text{Enc}(\mathbf{pk}, m_b)$ 
5:  $\hat{b} \xleftarrow{\$} \mathcal{A}(1^\lambda, \mathbf{pk}, c^*)$ 
6: return  $\llbracket b = \hat{b} \rrbracket$ 

```

Figure 4: The IND-CPA game

90 The *advantage* $\epsilon_{\text{IND-CPA}}$ of an IND-CPA adversary A is defined by

$$\text{Adv}_{\text{IND-CPA}}(A) = \left| P[\hat{b} = b] - \frac{1}{2} \right|$$

2.2 Key encapsulation mechanism

A key encapsulation mechanism KEM is a collection of three routines (**Gen**, **Encap**, **Decap**) defined over some ciphertext space \mathcal{C} and some key space \mathcal{K} . The key generation routine takes the security parameter 1^λ and outputs a keypair $(\mathbf{pk}, \mathbf{sk}) \xleftarrow{\$} \text{Gen}(1^\lambda)$. **Encap**(\mathbf{pk}) is a probabilistic routine that takes a public key \mathbf{pk} and outputs a pair of values (c, K) where $c \in \mathcal{C}$ is the encapsulation (or ciphertext) of the shared secret $k \in \mathcal{K}$. **Decap**(\mathbf{sk}, c) is a deterministic routine that takes the secret key \mathbf{sk} and the encapsulation c and returns the shared secret k if the ciphertext is valid, or the rejection symbol \perp if the ciphertext is invalid.

The IND-CCA security of a KEM is defined by an adversarial game in which an adversary's goal is to distinguish pseudorandom shared secret (generated by running the **Encap** routine) and a truly random value.

Algorithm 5 IND-CCA game for KEM

```

1:  $(\mathbf{pk}, \mathbf{sk}) \xleftarrow{\$} \text{Gen}(1^\lambda)$ 
2:  $(c^*, k_0) \xleftarrow{\$} \text{Encap}(\mathbf{pk})$ 
3:  $k_1 \xleftarrow{\$} \mathcal{K}$ 
4:  $b \xleftarrow{\$} \{0, 1\}$ 
5:  $\hat{b} \xleftarrow{\$} \mathcal{A}_{\text{IND-CCA}}^{\mathcal{O}^{\text{Decap}}}(1^\lambda, \mathbf{pk}, c^*, k_b)$ 
6: return  $\llbracket \hat{b} = b \rrbracket$ 

```

Figure 5: The KEM-IND-CCA2 game

The decapsulation oracle $\mathcal{O}^{\text{Decap}}$ takes a ciphertext c and returns the output of the **Decap** routine using the secret key. The advantage $\epsilon_{\text{IND-CCA}}$ of an IND-CCA adversary $\mathcal{A}_{\text{IND-CCA}}$ is defined by

$$\epsilon_{\text{IND-CCA}} = \left| P[\hat{b} = b] - \frac{1}{2} \right|$$

2.3 Message authentication code

A message authentication code MAC is a collection of routines (**Sign**, **Verify**) defined over some key space \mathcal{K} , some message space \mathcal{M} , and some tag space \mathcal{T} . The signing routine **Sign**(k, m) takes the secret key $k \in \mathcal{K}$ and some message, and outputs a tag t . The verification routine **Verify**(k, m, t) takes the triplet of secret key, message, and tag, and outputs 1 if the message-tag pair is valid under the secret key, or 0 otherwise.

The security of a MAC is defined in an adversarial game in which an adversary, with access to some signing oracle $\mathcal{O}_{\text{sign}}(m)$, tries to forge a new valid message-tag pair that has never been queried before. The existential unforgeability under chosen message attack (EUF-CMA) game is shown below:

Algorithm 6 The EUF-CMA game

```

1:  $k^* \xleftarrow{\$} \mathcal{K}$ 
2:  $(\hat{m}, \hat{t}) \xleftarrow{\$} \mathcal{A}^{\mathcal{O}_{\text{sign}}}()$ 
3: return  $\llbracket \text{Verify}(k^*, \hat{m}, \hat{t}) \wedge (\hat{m}, \hat{t}) \notin \mathcal{O}_{\text{sign}} \rrbracket$ 

```

Figure 6: The EUF-CMA game

116 The advantage $\epsilon_{\text{EUF-CMA}}$ of the existential forgery adversary is the probability that it
 117 wins the EUF-CMA game.

2.4 Modular Fujisaki-Okamoto transformation

119 The Fujisaki-Okamoto transformation (FOT) [FO99] is a generic transformation that
 120 takes a PKE with weaker security (such as OW-CPA or IND-CPA) and outputs a PKE with
 121 stronger security. A later variation [HHK17] improved the original construction in [FO99]
 122 by accounting for decryption failures, tightening security bounds, and providing a modular
 123 construction that first transforms OW-CPA/IND-CPA PKE into OW-PCVA PKE by providing
 124 ciphertext integrity through re-encryption (the T transformation), then transforming the
 125 OW-PCVA PKE into an IND-CCA KEM (the U transformation).

126 Particularly relevant to our results are two variations of the U transformation: U^\perp
 127 (KEM with explicit rejection) and U^\perp (KEM with implicit rejection). If PKE is OW-PCVA
 128 secure, then U^\perp transforms PKE into an IND-CCA secure KEM^\perp :

129 **Theorem 1.** *For any IND-CCA adversary \mathcal{A}_{KEM} against KEM^\perp with advantage ϵ_{KEM} issuing*
 130 *at most q_D decapsulation queries and at most q_H hash queries, there exists an OW-PCVA*
 131 *adversary \mathcal{A}_{PKE} against the underlying PKE with advantage ϵ_{PKE} that makes at most q_H*
 132 *queries to PCO and CVO such that*

$$\epsilon_{\text{KEM}} \leq \epsilon_{\text{PKE}}$$

133 Similarly, if PKE is OW-PCA secure, then U^\perp transforms PKE into an IND-CCA secure KEM^\perp

134 **Theorem 2.** *For any IND-CCA adversary \mathcal{A}_{KEM} against KEM^\perp with advantage ϵ_{KEM} issuing*
 135 *at most q_D decapsulation queries and at most q_H hash queries, there exists an OW-CPA*
 136 *adversary \mathcal{A}_{PKE} against the underlying PKE with advantage ϵ_{PKE} issuing at most q_H queries*
 137 *to PCO such that:*

$$\epsilon_{\text{KEM}} \leq \frac{q_H}{|\mathcal{M}_{\text{PKE}}|} + \epsilon_{\text{PKE}}$$

138 The modularity of the T and U transformation allows us to tweak only the T transfor-
 139 mation (see section 3), obtain OW-PCVA security, then automatically get IND-CCA security
 140 for free. This means that we can directly apply our contribution to existing KEM's already
 141 using this modular transformation, such as ML-KEM [KE23], and obtain performance
 142 improvements while maintaining comparable levels of security (see section 4.1).

3 The “encrypt-then-MAC” transformation

143 Let $\text{PKE}(\text{Gen}, \text{Enc}, \text{Dec})$ be a public-key encryption scheme. Let MAC be a deterministic
 144 message authentication code. Let $G : \mathcal{M}_{\text{PKE}} \rightarrow \mathcal{K}_{\text{MAC}}$ and $H : \{0, 1\}^* \rightarrow \mathcal{K}_{\text{KEM}}$ be hash
 145 functions, where \mathcal{K}_{KEM} denote the set of all possible session keys. The EtM transformation
 146

147 outputs a key encapsulation mechanism $\text{KEM}_{\text{EtM}}(\text{Gen}_{\text{EtM}}, \text{Encap}_{\text{EtM}}, \text{Decap}_{\text{EtM}})$. The three
 148 routines are described in figure 7.

Algorithm 7 Gen_{EtM}

```

1:  $(\text{pk}, \text{sk}_{\text{PKE}}) \xleftarrow{\$} \text{Gen}(1^\lambda)$ 
2:  $z \xleftarrow{\$} \mathcal{M}_{\text{PKE}}$ 
3:  $\text{sk} \leftarrow (\text{sk}_{\text{PKE}}, z)$ 
4: return  $(\text{pk}, \text{sk})$ 
    
```

Algorithm 8 $\text{Encap}_{\text{EtM}}(\text{pk})$

```

1:  $m \xleftarrow{\$} \mathcal{M}_{\text{PKE}}$ 
2:  $k \leftarrow G(m)$ 
3:  $c_{\text{PKE}} \xleftarrow{\$} \text{Enc}(\text{pk}, m)$ 
4:  $t \leftarrow \text{MAC}(k, c_{\text{PKE}})$ 
5:  $K \leftarrow H(m, c_{\text{PKE}})$ 
6:  $c \leftarrow (c_{\text{PKE}}, K)$ 
7: return  $(c, K)$ 
    
```

Algorithm 9 $\text{Decap}_{\text{EtM}}(\text{sk}, c)$

```

1:  $(c_{\text{PKE}}, t) \leftarrow c$ 
2:  $(\text{sk}_{\text{PKE}}, z) \leftarrow \text{sk}$ 
3:  $\hat{m} \leftarrow \text{Dec}(\text{sk}_{\text{PKE}}, c_{\text{PKE}})$ 
4:  $\hat{k} \leftarrow G(\hat{m})$ 
5: if  $\text{MAC}(\hat{k}, c_{\text{PKE}}) \neq t$  then
6:   return  $H(z, c_{\text{PKE}})$ 
7: end if
8: return  $H(\hat{m}, c_{\text{PKE}})$ 
    
```

Figure 7: KEM_{EtM} routines

149 **Theorem 3.** For every *IND-CCA2* adversary A against KEM_{EtM} that makes q_D decapsulation
 150 queries, there exists an *OW-PCA* adversary B who makes at least q_D plaintext-checking
 151 queries against the underlying PKE such that

$$\text{Adv}_{\text{IND-CCA2}}(A) \leq q_D \cdot \epsilon_{\text{MAC}} + 2 \cdot \text{Adv}_{\text{OW-PCA}}(B)$$

152 *Proof.* We will prove using a sequence of games. The complete sequence of games is shown
 153 in figure 8

Algorithm 10 Sequence of games $G_0 - G_3$

```

1:  $(\text{pk}, \text{sk}) \xleftarrow{\$} \text{Gen}(1^\lambda)$ 
2:  $(m^*, z) \xleftarrow{\$} \mathcal{M}_{\text{PKE}}$ 
3:  $k^* \leftarrow G(m^*)$   $\triangleright G_0 - G_1$ 
4:  $k^* \xleftarrow{\$} \mathcal{K}_{\text{MAC}}$   $\triangleright G_2 - G_3$ 
5:  $c_{\text{PKE}}^* \xleftarrow{\$} \text{Enc}(\text{pk}, m^*)$ 
6:  $t^* \leftarrow \text{MAC}(k^*, c_{\text{PKE}}^*)$ 
7:  $c^* \leftarrow (c_{\text{PKE}}^*, t^*)$ 
8:  $K_0 \leftarrow H(m^*, c_{\text{PKE}}^*)$   $\triangleright G_0 - G_2$ 
9:  $K_0 \xleftarrow{\$} \mathcal{K}_{\text{KEM}}$   $\triangleright G_3$ 
10:  $K_1 \xleftarrow{\$} \mathcal{K}_{\text{KEM}}$ 
11:  $b \xleftarrow{\$} \{0, 1\}$ 
12:  $\hat{b} \leftarrow A^{\mathcal{O}^{\text{Decap}}}(1^\lambda, \text{pk}, c^*, K_b)$   $\triangleright G_0$ 
13:  $\hat{b} \leftarrow A^{\mathcal{O}_1^{\text{Decap}}}(1^\lambda, \text{pk}, c^*, K_b)$   $\triangleright G_1 - G_3$ 
14: return  $\llbracket \hat{b} = b \rrbracket$ 

```

Algorithm 11 $\mathcal{O}^{\text{Decap}}(c)$

```

1:  $(c_{\text{PKE}}, t) \leftarrow c$ 
2:  $\hat{m} \leftarrow \text{Dec}(\text{sk}_{\text{PKE}}, c_{\text{PKE}})$ 
3:  $\hat{k} \leftarrow G(\hat{m})$ 
4: if  $\text{MAC}(\hat{k}, c_{\text{PKE}}) = t$  then
5:   return  $H(\hat{m}, c_{\text{PKE}})$ 
6: end if
7: return  $H(z, c_{\text{PKE}})$ 

```

Algorithm 12 $\mathcal{O}_1^{\text{Decap}}(c)$

```

1:  $(c_{\text{PKE}}, t) \leftarrow c$ 
2: if  $\exists (\tilde{m}, \tilde{k}) \in \mathcal{L}^G : \text{Dec}(\text{sk}_{\text{PKE}}, c_{\text{PKE}}) = \tilde{m} \wedge \text{MAC}(\tilde{k}, c_{\text{PKE}}) = t$  then
3:   return  $H(\tilde{m}, c_{\text{PKE}})$ 
4: end if
5: return  $H(z, c_{\text{PKE}})$ 

```

Figure 8: Sequence of games, true decap oracle $\mathcal{O}^{\text{Decap}}$ and simulated oracle $\mathcal{O}_1^{\text{Decap}}$

154 *Game 0* is the standard IND-CCA2 game for a key encapsulation mechanism.

155 *Game 1* is identical to *Game 0* except for that the decapsulation oracle $\mathcal{O}^{\text{Decap}}$ (algorithm
156 11) is replaced with a simulated decapsulation oracle $\mathcal{O}_1^{\text{Decap}}$ (algorithm 12). If $\mathcal{O}_1^{\text{Decap}}$
157 accepts the queried ciphertext $c = (c_{\text{PKE}}, t)$ and outputs the true session key $K \leftarrow H(\tilde{m}, c_{\text{PKE}})$,
158 then the queried ciphertext must be honestly generated, which means that $\mathcal{O}^{\text{Decap}}$ must
159 also accept the queried ciphertext and output the true session key. If $\mathcal{O}^{\text{Decap}}$ rejects the
160 queried ciphertext $c = (c_{\text{PKE}}, t)$ and outputs the implicit rejection $K \leftarrow H(z, c_{\text{PKE}})$, then the
161 tag t is invalid under the MAC key $k \leftarrow G(\text{Dec}(\text{sk}_{\text{PKE}}, c_{\text{PKE}}))$. Since for a given ciphertext
162 c_{PKE} , the correct MAC key is fixed, there could not be a matching hash query (m, k) such

that m is the correct decryption and k can validate the incorrect tag. Therefore, $\mathcal{O}_1^{\text{Decap}}$ must also reject the queried ciphertext and output the implicit rejection.

This means that game 0 and game 1 differ when $\mathcal{O}^{\text{Decap}}$ accepts the queried ciphertext $c = (c_{\text{PKE}}, t)$ but $\mathcal{O}_1^{\text{Decap}}$ rejects it, which means that t is a valid tag for c_{PKE} under the correct MAC key $k \leftarrow G(\text{Dec}(\text{sk}_{\text{PKE}}, c_{\text{PKE}}))$ but such key is never queried by the adversary. Under the random oracle model, from the adversary's perspective, such k is an unknown and uniformly random key, so producing a valid tag under such key constitutes a forgery against the MAC. Denote the probability of forgery against unknown uniformly random MAC key by ϵ_{MAC} , then the probability that the two decapsulation oracles disagree on one or more queries is at most $q_D \cdot \epsilon_{\text{MAC}}$. Finally, by the difference lemma,

$$\text{Adv}_0(A) - \text{Adv}_1(A) \leq q_D \cdot \epsilon_{\text{MAC}}$$

Note that ϵ_{MAC} quantifies the probability that an adversary can produce forgery for a unknown key without access to a signing oracles. While this is not a standard security definition for MAC, this probability is straightforward to estimate for some classes of MACs. As will be discussed in section 4.1, with a Carter-Wegman-like one-time MAC instantiated with message length L and a finite field with F elements, such probability is at most $\epsilon_{\text{MAC}} \leq \frac{L+1}{F}$.

Game 2 is identical to *Game 1*, except for that when the challenger generates the challenge ciphertext $c^* = (c_{\text{PKE}}^*, t^*)$, the tag t^* is computed using a uniformly random key $k^* \leftarrow \mathcal{K}_{\text{MAC}}$ instead of a pseudorandom key derived from hashing the challenge plaintext.

Under the random oracle model, game 2 and game 1 are statistically identical to the adversary A , unless A queries G with m^* . Denote the probability that A queries G with m^* by $P[\text{QUERY } G^*]$, then:

$$\text{Adv}_1(A) - \text{Adv}_2(A) \leq P[\text{QUERY } G^*]$$

Game 3 is identical to *Game 2*, except for that K_0 is a uniformly random session key instead of a pseudorandom session key derived from the challenge plaintext-ciphertext pair. Under the random oracle model, game 3 and game 2 are statistically identical unless the adversary A queries H with (m^*, \cdot) . Denote the probability that A makes such H query by $P[\text{QUERY } H^*]$, then:

$$\text{Adv}_2(A) - \text{Adv}_3(A) \leq P[\text{QUERY } H^*]$$

In game 3, both K_0 and K_1 are uniformly random. There is no statistical difference between the two session keys, so no adversary can have any advantage: $\text{Adv}_3(A) = 0$.

Now consider an OW-PCA adversary B simulating game 3 for A :

1. When B receives its public key pk , B passes pk to A
2. B can sample the implicit rejection z by itself
3. B can simulate both hash oracles G and H for A
4. B can simulate $\mathcal{O}_1^{\text{Decap}}$ for A . Instead of checking if $\text{Dec}(\text{sk}_{\text{PKE}}, c) = \tilde{m}$, B can use its access to the plaintext-checking oracle and check if $\text{PCO}(\tilde{m}, c) = 1$. This means that for every decapsulation query B services, B needs to make at least one plaintext-checking query. Therefore, B needs to make at least q_D plaintext-checking query.
5. When B receives its challenge ciphertext c_{PKE}^* , it can sample a uniformly random key $k^* \xleftarrow{\$} \mathcal{K}_{\text{MAC}}$, produce the corresponding tag $t^* \leftarrow \text{MAC}(k^*, c_{\text{PKE}}^*)$, and sample a uniformly random session keys $K_0, K_1 \xleftarrow{\$} \mathcal{K}_{\text{KEM}}$. B then passes $c^* = (c_{\text{PKE}}^*, t^*)$ as the challenge ciphertext and a coin-flip K_b as the challenge session key.

If A ever queries G or H with the decryption of c_{PKE}^* , B will be able to detect it using the plaintext-checking oracle. From where B is guaranteed to win the OW-PCA game. Therefore:

$$\begin{aligned} P[\text{QUERY } G^*] &\leq \text{Adv}_{\text{OW-PCA}}(B) \\ P[\text{QUERY } H^*] &\leq \text{Adv}_{\text{OW-PCA}}(B) \end{aligned}$$

Combining all inequalities above, we have:

$$\text{Adv}_0(A) \leq q_D \cdot \epsilon_{\text{MAC}} + 2\text{Adv}_{\text{OW-PCA}}(B)$$

□

4 Application to Kyber

CRYSTALS-Kyber [BDK⁺18][ABD⁺19] is an IND-CCA2 secure key encapsulation mechanism that first constructs an IND-CPA secure public key encryption scheme whose security is based on the Module Learning with Error Problem (MLWE), then applies a generic transformation using *de-randomization* and *re-encryption* [HHK17]. The use of *re-encryption* for providing rigidity means that the decapsulation routine needs to run the encryption routine for verifying ciphertext integrity. Unfortunately for Kyber, the encryption routine consumes both more CPU cycles and carries a larger memory footprint. Therefore, applying an alternate transformation that replaces *re-encryption* with a MAC verification will bring substantial performance enhancement. This is particularly applicable to communication protocols such as TLS 1.3, where clients, often constrained environments, need to run the decapsulation routine.

The IND-CPA PKE of Kyber (algorithms 4, 5, 6 in [ABD⁺19]) is not OW-PCA secure. A plaintext-checking attack [RRCB19] can recover the secret key using a few maliciously constructed plaintext-checking queries for each coefficient of a Kyber secret key (there are 512, 768, and 1024 coefficients in the secret key depending on the desired security level). However, we propose mitigation for plaintext-checking attack at protocol level by requiring each keypair to be used for decryption only once. If a decryption fails, then the secret key should be discarded and the key exchange terminated, and a new keypair should be generated to restart the key exchange. In fact, such ephemeral key exchange is already required in TLS 1.3 for forward secrecy, so requiring the one-time use of keypair does not introduce additional operational cost for protocols such as TLS 1.3.

The routines of CCAKEM can be found in Algorithm 7, 8, 9 in [ABD⁺19]. We modify Algorithms 8 and 9 using authenticated encryption (AE) mode as follows where Algorithm 7 is unchanged.

Algorithm 13 Kyber.CCAKEM.KeyGen()

- 1: $z \xleftarrow{\$} \mathcal{B}^{32}$ ▷ Randomly sample 32 bytes (256 bits)
 - 2: $(\text{pk}, \text{sk}') \xleftarrow{\$} \text{Kyber.CPAPKE.KeyGen}()$
 - 3: $\text{sk} = (\text{sk}', \text{pk}, H(\text{pk}), z)$ ▷ H is instantiated with SHA3-256
 - 4: **return** (pk, sk)
-

Algorithm 14 $\text{Kyber.CCAKEM.Encap}^+(\text{pk})$

```

1:  $m \xleftarrow{\$} \mathcal{B}^{32}$  ▷ Do not output system RNG directly
2:  $(\bar{K}, r, K_{\text{mac}}) = G(m' \| H(\text{pk}))$  ▷ G is instantiated with SHA3-512
3:  $c \leftarrow \text{Kyber.CPAPKE.Enc}(\text{pk}, m', r)$  ▷ Because  $r$  is set, CPAPKE is deterministic
4:  $t_1 = \text{MAC}(K_{\text{mac}}, c)$ 
5:  $K = \text{KDF}(\bar{K} \| t_1)$  ▷ KDF is instantiated with Shake256
6:  $c \leftarrow (c, t_1)$ 
7: return  $(c, K)$ 

```

Algorithm 15 $\text{Kyber.CCAKEM.Decap}^+(\text{sk}, c, t_1, t_2)$ **Require:** Secret key $\text{sk} = (\text{sk}', \text{pk}, H(\text{pk}), z)$ **Require:** Kyber.CPAPKE Ciphertext c and Tags t_1, t_2

```

1:  $(\text{sk}', \text{pk}, h, z) \leftarrow \text{sk}$  ▷ Unpack the secret key;  $h$  is the hash of  $\text{pk}$ 
2:  $(c, t_1) \leftarrow c$ 
3:  $m = \text{Kyber.CPAPKE.Dec}(\text{sk}', c)$ 
4:  $(\bar{K}, r, K_{\text{mac}}) = G(m' \| h)$ 
5:  $t'_1 = \text{MAC}(K_{\text{mac}}, c)$ 
6: if  $t'_1 = t_1$  then
7:    $K = \text{KDF}(\bar{K} \| t_1)$ 
8:   return  $K$ 
9: else
10:   Abort
11: end if

```

Remark 1. If c is manipulated, then the verification of t_1 will be failed. In this case, there is no K outputted from the decap^+ . So the attacks described in the following subsections won't work. We also added the key confirmation, which is tag t_2 .

Note for authenticated encryption, tags t_i 's are necessary for the inputs.

In fact, this is authenticated encryption instead of EtM. So we should change that, called authenticated encryption. Please do not change my notation. They have their meanings in AE mode.

4.1 MAC performance

When instantiating an instance $\text{KEM}(\text{Gen}, \text{Encap}, \text{Decap})$, there are a variety of possible MAC's to choose from. For each of the security level and choice of MAC, we measured the number of CPU cycles needed to produce a digest of the unauthenticated ciphertext. The median (top) and mean (bottom) measurements are reported in table 1.

Table 1: Standalone MAC performances

Name	Security	measurement	768 bytes	1088 bytes	1568 bytes
CMAC	many-time	median	5022	5442	6090
		mean	5131	5578	6154
KMAC-256	many-time	median	7934	9862	11742
		mean	8594	10693	12319
GMAC	one-time	median	2778	2756	2762
		mean	2843	2780	2919
Poly1305	one-time	median	1128	1218	1338
		mean	1435	1504	1625

Based on the security reduction in section 3 we chose standalone GMAC and Poly1305 for their substantial performance advantage over other constructions such as CBC-MAC and KMAC. We then modified the reference implementation (TODO: citation needed <https://github.com/pq-crystals/kyber>) according to algorithms 13, 14, and 15. We measured the median and mean CPU cycles needed to run each of the routines. The measurements are listed in table 2

Table 2: Kyber-AE performance measurements

Name	Security level	measurement	Gen	Encap	Decap
Kyber512	128 bits	median			
		mean			
Kyber768	192 bits	median			
		mean			
Kyber1024	256 bits	median			
		mean			
KyberAE512 w/ GMAC	128 bits	median			
		mean			
KyberAE768 w/ GMAC	192 bits	median			
		mean			
KyberAE1024 w/ GMAC	256 bits	median			
		mean			
KyberAE512 w/ Poly1305	128 bits	median			
		mean			
KyberAE768 w/ Poly1305	192 bits	median			
		mean			
KyberAE1024 w/ Poly1305	256 bits	median			
		mean			

5 Conclusions and future works

References

- [ABD⁺19] Roberto Avanzi, Joppe Bos, Léo Ducas, Eike Kiltz, Tancrede Lepoint, Vadim Lyubashevsky, John M Schanck, Peter Schwabe, Gregor Seiler, and Damien Stehlé. Crystals-kyber algorithm specifications and supporting documentation. *NIST PQC Round*, 2(4):1–43, 2019.
- [BDK⁺18] Joppe Bos, Léo Ducas, Eike Kiltz, Tancrede Lepoint, Vadim Lyubashevsky, John M Schanck, Peter Schwabe, Gregor Seiler, and Damien Stehlé. Crystals-kyber: a cca-secure module-lattice-based kem. In *2018 IEEE European Symposium on Security and Privacy (EuroS&P)*, pages 353–367. IEEE, 2018.
- [BN00] Mihir Bellare and Chanathip Namprempre. Authenticated encryption: Relations among notions and analysis of the generic composition paradigm. In *International Conference on the Theory and Application of Cryptology and Information Security*, pages 531–545. Springer, 2000.
- [BP18] Daniel J Bernstein and Edoardo Persichetti. Towards kem unification. *Cryptography ePrint Archive*, 2018.
- [FO99] Eiichiro Fujisaki and Tatsuaki Okamoto. Secure integration of asymmetric and symmetric encryption schemes. In *Annual international cryptology conference*, pages 537–554. Springer, 1999.

- 271 [HHK17] Dennis Hofheinz, Kathrin Hövelmanns, and Eike Kiltz. A modular analysis of
272 the fujisaki-okamoto transformation. In *Theory of Cryptography Conference*,
273 pages 341–371. Springer, 2017.
- 274 [KE23] NIST Module-Lattice-Based Key-Encapsulation. Mechanism standard. *NIST*
275 *Post-Quantum Cryptography Standardization Process; NIST: Gaithersburg, MD,*
276 *USA*, 2023.
- 277 [RRCB19] Prasanna Ravi, Sujoy Sinha Roy, Anupam Chattopadhyay, and Shivam Bhasin.
278 Generic side-channel attacks on cca-secure lattice-based pke and kem schemes.
279 *Cryptology ePrint Archive*, 2019.
- 280 [Sho04] Victor Shoup. Sequences of games: a tool for taming complexity in security
281 proofs. *cryptology eprint archive*, 2004.