

Fast Fujisaki-Okamoto transformation using encrypt-then-mac and applications to Kyber

Anonymous Submission

Abstract. The modular Fujisaki-Okamoto (FO) transformation takes public-key encryption with weaker security and constructs a key encapsulation mechanism (KEM) with indistinguishability under adaptive chosen ciphertext attacks. While the modular FO transform enjoys tight security bound and quantum resistance, it also suffers from computational inefficiency and vulnerabilities to side-channel attacks due to using de-randomization and re-encryption for providing ciphertext integrity. In this work, we propose an alternative KEM construction that achieves ciphertext integrity using a message authentication code (MAC) and instantiate a concrete instance using ML-KEM. Our experimental results showed that where the encryption routine incurs heavy computational cost, replacing re-encryption with MAC provides substantial performance improvements at comparable security level.

Keywords: Key encapsulation mechanism, post-quantum cryptography, lattice cryptography, Fujisaki-Okamoto transformation

1 Introduction

The Fujisaki-Okamoto transformation [FO99] is a generic construction that takes cryptographic primitives of lesser security and constructs a public-key encryption scheme with indistinguishability under adaptive chosen ciphertext attacks. Later works [HHK17] extended the original transformation to the construction of key encapsulation mechanism, which has been adopted by many post-quantum schemes such as Kyber [BDK⁺18], FrodoKEM [BCD⁺16], and SABER [DKSRV18].

The current state of the FO transformation enjoys tight security bound and quantum resistance [HHK17], but also leaves many deficiencies to be improved on. One such shortcoming is the use of re-encryption for providing ciphertext integrity [BP18], which requires the decapsulation routine to run the encryption routine as a subroutine. In many post-quantum schemes, such as Kyber, the encryption routine is substantially more expensive than the decryption routine, so running the encryption routine in the decapsulation routine inflates the computational cost of the decapsulator. In addition, running the encryption as a subroutine introduces risks of side-channel vulnerabilities that may expose the plaintext or the secret key [RRCB19][UXT⁺22].

The problem of ciphertext integrity was solved in symmetric cryptography: given a semantically secure symmetric cipher and an existentially unforgeable message authentication code, combining them using “encrypt-then-mac” provides authenticated encryption [BN00]. We took inspiration from this strategy and applied a similar technique to transform an IND-CPA secure public-key encryption scheme into an IND-CCA2 secure key encapsulation mechanism. Using a message authentication code for ciphertext integrity replaces the re-encryption step in decryption with the computation of an authenticator, which offers significant performance improvements while maintaining comparable level of security.

The main challenge in applying “encrypt-then-mac” to public-key cryptography is the lack of a pre-shared symmetric key. We proposed to derive the symmetric key by hashing the plaintext message. In section 3, we prove that under the random oracle model, if the

input public-key encryption scheme is one-way secure against plaintext-checking attack and the input message authentication code is one-time existentially unforgeable, then the transformed key encapsulation mechanism is IND-CCA2 secure.

In section 4, we instantiate concrete instances of our constructions by combining Kyber with GMAC and Poly1305. Our experimental results showed that replacing re-encryption with computing authenticator leads to significant performance improvements in the decapsulation routine while incurring only minimal overhead in the encapsulation routine.

2 Preliminaries and previous results

2.1 Public-key encryption scheme

A public key encryption scheme PKE is a collection of three routines (KeyGen , Enc , Dec) defined over some message space \mathcal{M} and some ciphertext space \mathcal{C} . Where the encryption routine is probabilistic, the source of randomness is denoted by the coin space \mathcal{R} .

The encryption routine $\text{Enc}(\text{pk}, m)$ takes a public key, a plaintext message, and outputs a ciphertext $c \in \mathcal{C}$. Where the encryption routine is probabilistic, specifying a pseudorandom seed $r \in \mathcal{R}$ will make the encryption routine behave deterministically. The decryption routine $\text{Dec}(\text{sk}, c)$ takes a secret key, a ciphertext, and outputs the decryption \hat{m} if the ciphertext is valid under the given secret key, or the rejection symbol \perp if the ciphertext is invalid.

We discuss the security of a PKE using the sequence of games described in [Sho04]. Specifically, we first define the OW-ATK as they pertain to a public key encryption scheme. In later section we will define the IND-CCA game as it pertains to a key encapsulation mechanism.

Algorithm 1 The OW-ATK game

```

1:  $(\text{pk}, \text{sk}) \xleftarrow{\$} \text{KeyGen}(1^\lambda)$ 
2:  $m^* \xleftarrow{\$} \mathcal{M}$ 
3:  $c^* \xleftarrow{\$} \text{Enc}(\text{pk}, m^*)$ 
4:  $\hat{m} \xleftarrow{\$} \mathcal{A}^{\mathcal{O}_{\text{ATK}}}(1^\lambda, \text{pk}, c^*)$ 
5: return  $\llbracket m^* = \hat{m} \rrbracket$ 

```

Figure 1: The OW-ATK game

Algorithm 2 PCO($m \in \mathcal{M}, c \in \mathcal{C}$)

```

1: return  $\llbracket \text{Dec}(\text{sk}, c) = m \rrbracket$ 

```

Figure 2: Plaintext-checking oracle

In the OW-ATK game (see figure 1), an adversary's goal is to recover the decryption of a randomly generated ciphertext. A challenger randomly samples a keypair and a challenge plaintext m^* , encrypts the challenge plaintext $c^* \xleftarrow{\$} \text{Enc}(\text{pk}, m^*)$, then gives pk and c^* to the adversary A . The adversary A , with access to some oracle \mathcal{O}_{ATK} , outputs a guess decryption \hat{m} . A wins the game if its guess \hat{m} is equal to the challenge plaintext m^* . The *advantage* $\text{Adv}_{\text{OW-ATK}}$ of an adversary in this game is the probability that it wins the game:

$$\text{Adv}_{\text{OW-ATK}}(A) = P \left[A(\text{pk}, c^*) = m^* \mid (\text{pk}, \text{sk}) \xleftarrow{\$} \text{KeyGen}(); m^* \xleftarrow{\$} \mathcal{M}; c^* \xleftarrow{\$} \text{Enc}(\text{pk}, m^*) \right]$$

The capabilities of the oracle \mathcal{O}_{ATK} depends on the choice of security goal ATK. Particularly relevant to our result is security against plaintext-checking attack (PCA), for which the adversary has access to a plaintext-checking oracle (PCO) (see figure 2). A PCO takes

as input a plaintext-ciphertext pair (m, c) and returns **True** if m is the decryption of c or **False** otherwise.

2.2 Key encapsulation mechanism (KEM)

A key encapsulation mechanism is a collection of three routines (**KeyGen**, **Encap**, **Decap**) defined over some ciphertext space \mathcal{C} and some key space \mathcal{K} . The key generation routine takes the security parameter 1^λ and outputs a keypair $(\mathbf{pk}, \mathbf{sk}) \xleftarrow{\$} \text{KeyGen}(1^\lambda)$. **Encap**(\mathbf{pk}) is a probabilistic routine that takes a public key \mathbf{pk} and outputs a pair of values (c, K) where $c \in \mathcal{C}$ is the ciphertext (also called encapsulation) and $K \in \mathcal{K}$ is the shared secret (also called session key). **Decap**(\mathbf{sk}, c) is a deterministic routine that takes the secret key \mathbf{sk} and the encapsulation c and returns the shared secret K if the ciphertext is valid. Some KEM constructions use explicit rejection, where if c is invalid then **Decap** will return a rejection symbol \perp ; other KEM constructions use implicit rejection, where if c is invalid then **Decap** will return a “fake” session key that depends on the ciphertext and some other secret values.

The IND-CCA security of a KEM is defined by an adversarial game in which an adversary’s goal is to distinguish pseudorandom shared secret (generated by running the **Encap** routine) and a truly random value.

Algorithm 3 IND-CCA game for KEM

```

1:  $(\mathbf{pk}, \mathbf{sk}) \xleftarrow{\$} \text{KeyGen}(1^\lambda)$ 
2:  $(c^*, K_0) \xleftarrow{\$} \text{Encap}(\mathbf{pk})$ 
3:  $K_1 \xleftarrow{\$} \mathcal{K}$ 
4:  $b \xleftarrow{\$} \{0, 1\}$ 
5:  $\hat{b} \xleftarrow{\$} A^{\mathcal{O}_{\text{Decap}}}(1^\lambda, \mathbf{pk}, c^*, K_b)$ 
6: return  $\llbracket \hat{b} = b \rrbracket$ 
```

Algorithm 4 $\mathcal{O}_{\text{Decap}}(c)$

```

1: return Decap( $\mathbf{sk}, c$ )
```

Figure 3: The KEM-IND-CCA2 game

Figure 4: Decapsulation oracle

The decapsulation oracle $\mathcal{O}^{\text{Decap}}$ takes a ciphertext c and returns the output of the **Decap** routine using the secret key. The advantage $\epsilon_{\text{IND-CCA}}$ of an IND-CCA adversary $\mathcal{A}_{\text{IND-CCA}}$ is defined by

$$\text{Adv}_{\text{IND-CCA}}(A) = \left| P[A^{\mathcal{O}_{\text{Decap}}}(1^\lambda, \mathbf{pk}, c^*, K_b) = b] - \frac{1}{2} \right|$$

2.3 Message authentication code (MAC)

A message authentication code MAC is a collection of routines (**Sign**, **Verify**) defined over some key space \mathcal{K} , some message space \mathcal{M} , and some tag space \mathcal{T} . The signing routine **Sign**(k, m) takes the secret key $k \in \mathcal{K}$ and some message, and outputs a tag t . The verification routine **Verify**(k, m, t) takes the triplet of secret key, message, and tag, and outputs 1 if the message-tag pair is valid under the secret key, or 0 otherwise. Many MAC constructions are deterministic. For these constructions it is simpler to denote the signing routine by $t \leftarrow \text{MAC}(k, m)$ and perform verification using a simple comparison.

The security of a MAC is defined in an adversarial game in which an adversary, with access to some signing oracle $\mathcal{O}_{\text{Sign}}(m)$, tries to forge a new valid message-tag pair that

106 has never been queried before. The existential unforgeability under chosen message attack
 107 (EUF-CMA) game is shown below:

Algorithm 5 The EUF-CMA game

```

1:  $k^* \xleftarrow{\$} \mathcal{K}$ 
2:  $(\hat{m}, \hat{t}) \xleftarrow{\$} \mathcal{A}^{\mathcal{O}_{\text{sign}}}()$ 
3: return  $\llbracket \text{Verify}(k^*, \hat{m}, \hat{t}) \wedge (\hat{m}, \hat{t}) \notin \mathcal{O}_{\text{sign}} \rrbracket$ 

```

Figure 5: The EUF-CMA game

108 The advantage $\text{Adv}_{\text{EUF-CMA}}$ of the existential forgery adversary is the probability that it
 109 wins the EUF-CMA game.

110 We are also interested in one-time MAC, whose security goals are identical to normal
 111 MAC, except the signing oracle will only answer one signing query. In practice this means
 112 that the adversary only gets to see one message-tag pair, and the owner of the secret key
 113 only signs one message. One way to construct an unforgeable one-time MAC is to use a
 114 universal hash function (UHF), which computes the digest using finite field arithmetic.
 115 Each instance of the one-time MAC is parameterized by some maximal message length
 116 L and the finite field \mathbb{F} in which arithmetics are performed. The tag is computed by
 117 evaluating a polynomial whose coefficients are derived from the message:

$$H(k_1, (m_1, m_2, \dots, m_l)) = k_1^{l+1} + k_1^l m_1 + k_1^{l-1} m_2 + \dots + k_1 m_l$$

$$\text{MAC}((k_1, k_2), m) = H(k_1, m) + k_2$$

118 We claim without proof (though proof can be found in [BS20]) that the MAC construc-
 119 tion above is one-time existentially unforgeable.

120 2.4 Related works

121 The Fujisaki-Okamoto transformation (FOT) [FO99] is a generic transformation that
 122 takes a PKE with weaker security (such as OW-CPA or IND-CPA) and outputs a PKE with
 123 stronger security. A later variation [HHK17] improved the original construction in [FO99]
 124 by accounting for decryption failures, tightening security bounds, and providing a modular
 125 construction that first transforms OW-CPA/IND-CPA PKE into OW-PCVA PKE by providing
 126 ciphertext integrity through re-encryption (the T transformation), then transforming the
 127 OW-PCVA PKE into an IND-CCA KEM (the U transformation).

128 Particularly relevant to our results are two variations of the U transformation: U^\perp
 129 (KEM with explicit rejection) and U^χ (KEM with implicit rejection). If PKE is OW-PCVA
 130 secure, then U^\perp transforms PKE into an IND-CCA secure KEM^\perp :

131 **Theorem 1.** *For any IND-CCA adversary \mathcal{A}_{KEM} against KEM^\perp with advantage ϵ_{KEM} issuing*
 132 *at most q_D decapsulation queries and at most q_H hash queries, there exists an OW-PCVA*
 133 *adversary \mathcal{A}_{PKE} against the underlying PKE with advantage ϵ_{PKE} that makes at most q_H*
 134 *queries to PCO and CVO such that*

$$\epsilon_{\text{KEM}} \leq \epsilon_{\text{PKE}}$$

135 Similarly, if PKE is OW-PCA secure, then U^χ transforms PKE into an IND-CCA secure KEM^χ

136 **Theorem 2.** *For any IND-CCA adversary \mathcal{A}_{KEM} against KEM^χ with advantage ϵ_{KEM} issuing*
 137 *at most q_D decapsulation queries and at most q_H hash queries, there exists an OW-CPA*

adversary \mathcal{A}_{PKE} against the underlying PKE with advantage ϵ_{PKE} issuing at most q_H queries to PCO such that:

$$\epsilon_{KEM} \leq \frac{q_H}{|\mathcal{M}_{PKE}|} + \epsilon_{PKE}$$

The modularity of the T and U transformation allows us to tweak only the T transformation (see section 3), obtain OW-PCVA security, then automatically get IND-CCA security for free. This means that we can directly apply our contribution to existing KEM's already using this modular transformation, such as ML-KEM [KE23], and obtain performance improvements while maintaining comparable levels of security (see section 4.1).

3 The “encrypt-then-MAC” transformation

Let $PKE(\text{KeyGen}, \text{Enc}, \text{Dec})$ be a public-key encryption scheme. Let MAC be a deterministic message authentication code. Let $G : \mathcal{M}_{PKE} \rightarrow \mathcal{K}_{MAC}$ and $H : \{0, 1\}^* \rightarrow \mathcal{K}_{KEM}$ be hash functions, where \mathcal{K}_{KEM} denote the set of all possible session keys. The EtM transformation outputs a key encapsulation mechanism $KEM_{\text{EtM}}(\text{KeyGen}_{\text{EtM}}, \text{Encap}_{\text{EtM}}, \text{Decap}_{\text{EtM}})$. The three routines are described in figure 6.

Algorithm 6 $\text{KeyGen}_{\text{EtM}}$

```

1:  $(\text{pk}, \text{sk}_{\text{PKE}}) \xleftarrow{\$} \text{KeyGen}(1^\lambda)$ 
2:  $z \xleftarrow{\$} \mathcal{M}_{\text{PKE}}$ 
3:  $\text{sk} \leftarrow (\text{sk}_{\text{PKE}}, z)$ 
4: return  $(\text{pk}, \text{sk})$ 

```

Algorithm 7 $\text{Encap}_{\text{EtM}}(\text{pk})$

```

1:  $m \xleftarrow{\$} \mathcal{M}_{\text{PKE}}$ 
2:  $k \leftarrow G(m)$ 
3:  $c_{\text{PKE}} \xleftarrow{\$} \text{Enc}(\text{pk}, m)$ 
4:  $t \leftarrow \text{MAC}(k, c_{\text{PKE}})$ 
5:  $K \leftarrow H(m, c_{\text{PKE}})$ 
6:  $c \leftarrow (c_{\text{PKE}}, K)$ 
7: return  $(c, K)$ 

```

Algorithm 8 $\text{Decap}_{\text{EtM}}(\text{sk}, c)$

```

1:  $(c_{\text{PKE}}, t) \leftarrow c$ 
2:  $(\text{sk}_{\text{PKE}}, z) \leftarrow \text{sk}$ 
3:  $\hat{m} \leftarrow \text{Dec}(\text{sk}_{\text{PKE}}, c_{\text{PKE}})$ 
4:  $\hat{k} \leftarrow G(\hat{m})$ 
5: if  $\text{MAC}(\hat{k}, c_{\text{PKE}}) \neq t$  then
6:   return  $H(z, c_{\text{PKE}})$ 
7: end if
8: return  $H(\hat{m}, c_{\text{PKE}})$ 

```

Figure 6: KEM_{EtM} routines

151 **Theorem 3.** For every IND-CCA2 adversary A against KEM_{EtM} that makes q_D decapsulation
152 queries, there exists an OW-PCA adversary B who makes at least q_D plaintext-checking
153 queries against the underlying PKE such that

$$\text{Adv}_{\text{IND-CCA2}}(A) \leq q_D \cdot \epsilon_{\text{MAC}} + 2 \cdot \text{Adv}_{\text{OW-PCA}}(B)$$

154 *Proof.* We will prove using a sequence of games. The complete sequence of games is shown
155 in figure 7

Algorithm 9 Sequence of games $G_0 - G_3$

```

1:  $(\text{pk}, \text{sk}) \xleftarrow{\$} \text{KeyGen}(1^\lambda)$ 
2:  $(m^*, z) \xleftarrow{\$} \mathcal{M}_{\text{PKE}}$ 
3:  $k^* \leftarrow G(m^*)$   $\triangleright G_0 - G_1$ 
4:  $k^* \xleftarrow{\$} \mathcal{K}_{\text{MAC}}$   $\triangleright G_2 - G_3$ 
5:  $c_{\text{PKE}}^* \xleftarrow{\$} \text{Enc}(\text{pk}, m^*)$ 
6:  $t^* \leftarrow \text{MAC}(k^*, c_{\text{PKE}}^*)$ 
7:  $c^* \leftarrow (c_{\text{PKE}}^*, t^*)$ 
8:  $K_0 \leftarrow H(m^*, c_{\text{PKE}}^*)$   $\triangleright G_0 - G_2$ 
9:  $K_0 \xleftarrow{\$} \mathcal{K}_{\text{KEM}}$   $\triangleright G_3$ 
10:  $K_1 \xleftarrow{\$} \mathcal{K}_{\text{KEM}}$ 
11:  $b \xleftarrow{\$} \{0, 1\}$ 
12:  $\hat{b} \leftarrow A^{\mathcal{O}^{\text{Decap}}}(1^\lambda, \text{pk}, c^*, K_b)$   $\triangleright G_0$ 
13:  $\hat{b} \leftarrow A^{\mathcal{O}_1^{\text{Decap}}}(1^\lambda, \text{pk}, c^*, K_b)$   $\triangleright G_1 - G_3$ 
14: return  $\llbracket \hat{b} = b \rrbracket$ 

```

Algorithm 10 $\mathcal{O}^{\text{Decap}}(c)$

```

1:  $(c_{\text{PKE}}, t) \leftarrow c$ 
2:  $\hat{m} \leftarrow \text{Dec}(\text{sk}_{\text{PKE}}, c_{\text{PKE}})$ 
3:  $\hat{k} \leftarrow G(\hat{m})$ 
4: if  $\text{MAC}(\hat{k}, c_{\text{PKE}}) = t$  then
5:   return  $H(\hat{m}, c_{\text{PKE}})$ 
6: end if
7: return  $H(z, c_{\text{PKE}})$ 

```

Algorithm 11 $\mathcal{O}_1^{\text{Decap}}(c)$

```

1:  $(c_{\text{PKE}}, t) \leftarrow c$ 
2: if  $\exists (\tilde{m}, \tilde{k}) \in \mathcal{L}^G : \text{Dec}(\text{sk}_{\text{PKE}}, c_{\text{PKE}}) = \tilde{m} \wedge \text{MAC}(\tilde{k}, c_{\text{PKE}}) = t$  then
3:   return  $H(\tilde{m}, c_{\text{PKE}})$ 
4: end if
5: return  $H(z, c_{\text{PKE}})$ 

```

Figure 7: Sequence of games, true decap oracle $\mathcal{O}^{\text{Decap}}$ and simulated oracle $\mathcal{O}_1^{\text{Decap}}$

156 *Game 0* is the standard IND-CCA2 game for a key encapsulation mechanism.

157 *Game 1* is identical to *Game 0* except for that the decapsulation oracle $\mathcal{O}^{\text{Decap}}$ (algorithm
158 10) is replaced with a simulated decapsulation oracle $\mathcal{O}_1^{\text{Decap}}$ (algorithm 11). If $\mathcal{O}_1^{\text{Decap}}$
159 accepts the queried ciphertext $c = (c_{\text{PKE}}, t)$ and outputs the true session key $K \leftarrow H(\tilde{m}, c_{\text{PKE}})$,
160 then the queried ciphertext must be honestly generated, which means that $\mathcal{O}^{\text{Decap}}$ must
161 also accept the queried ciphertext and output the true session key. If $\mathcal{O}^{\text{Decap}}$ rejects the
162 queried ciphertext $c = (c_{\text{PKE}}, t)$ and outputs the implicit rejection $K \leftarrow H(z, c_{\text{PKE}})$, then the
163 tag t is invalid under the MAC key $k \leftarrow G(\text{Dec}(\text{sk}_{\text{PKE}}, c_{\text{PKE}}))$. Since for a given ciphertext
164 c_{PKE} , the correct MAC key is fixed, there could not be a matching hash query (m, k) such

that m is the correct decryption and k can validate the incorrect tag. Therefore, $\mathcal{O}_1^{\text{Decap}}$ must also reject the queried ciphertext and output the implicit rejection.

This means that game 0 and game 1 differ when $\mathcal{O}^{\text{Decap}}$ accepts the queried ciphertext $c = (c_{\text{PKE}}, t)$ but $\mathcal{O}_1^{\text{Decap}}$ rejects it, which means that t is a valid tag for c_{PKE} under the correct MAC key $k \leftarrow G(\text{Dec}(\text{sk}_{\text{PKE}}, c_{\text{PKE}}))$ but such key is never queried by the adversary. Under the random oracle model, from the adversary's perspective, such k is an unknown and uniformly random key, so producing a valid tag under such key constitutes a forgery against the MAC. Denote the probability of forgery against unknown uniformly random MAC key by ϵ_{MAC} , then the probability that the two decapsulation oracles disagree on one or more queries is at most $q_D \cdot \epsilon_{\text{MAC}}$. Finally, by the difference lemma,

$$\text{Adv}_0(A) - \text{Adv}_1(A) \leq q_D \cdot \epsilon_{\text{MAC}}$$

Note that ϵ_{MAC} quantifies the probability that an adversary can produce forgery for a unknown key without access to a signing oracles. While this is not a standard security definition for MAC, this probability is straightforward to estimate for some classes of MACs. As will be discussed in section 4.1, with a Carter-Wegman-like one-time MAC instantiated with message length L and a finite field with F elements, such probability is at most $\epsilon_{\text{MAC}} \leq \frac{L+1}{F}$.

Game 2 is identical to *Game 1*, except for that when the challenger generates the challenge ciphertext $c^* = (c_{\text{PKE}}^*, t^*)$, the tag t^* is computed using a uniformly random key $k^* \leftarrow \mathcal{K}_{\text{MAC}}$ instead of a pseudorandom key derived from hashing the challenge plaintext.

Under the random oracle model, game 2 and game 1 are statistically identical to the adversary A , unless A queries G with m^* . Denote the probability that A queries G with m^* by $P[\text{QUERY } G^*]$, then:

$$\text{Adv}_1(A) - \text{Adv}_2(A) \leq P[\text{QUERY } G^*]$$

Game 3 is identical to *Game 2*, except for that K_0 is a uniformly random session key instead of a pseudorandom session key derived from the challenge plaintext-ciphertext pair. Under the random oracle model, game 3 and game 2 are statistically identical unless the adversary A queries H with (m^*, \cdot) . Denote the probability that A makes such H query by $P[\text{QUERY } H^*]$, then:

$$\text{Adv}_2(A) - \text{Adv}_3(A) \leq P[\text{QUERY } H^*]$$

In game 3, both K_0 and K_1 are uniformly random. There is no statistical difference between the two session keys, so no adversary can have any advantage: $\text{Adv}_3(A) = 0$.

Now consider an OW-PCA adversary B simulating game 3 for A :

1. When B receives its public key pk , B passes pk to A
2. B can sample the implicit rejection z by itself
3. B can simulate both hash oracles G and H for A
4. B can simulate $\mathcal{O}_1^{\text{Decap}}$ for A . Instead of checking if $\text{Dec}(\text{sk}_{\text{PKE}}, c) = \tilde{m}$, B can use its access to the plaintext-checking oracle and check if $\text{PCO}(\tilde{m}, c) = 1$. This means that for every decapsulation query B services, B needs to make at least one plaintext-checking query. Therefore, B needs to make at least q_D plaintext-checking query.
5. When B receives its challenge ciphertext c_{PKE}^* , it can sample a uniformly random key $k^* \xleftarrow{\$} \mathcal{K}_{\text{MAC}}$, produce the corresponding tag $t^* \leftarrow \text{MAC}(k^*, c_{\text{PKE}}^*)$, and sample a uniformly random session keys $K_0, K_1 \xleftarrow{\$} \mathcal{K}_{\text{KEM}}$. B then passes $c^* = (c_{\text{PKE}}^*, t^*)$ as the challenge ciphertext and a coin-flip K_b as the challenge session key.

206 If A ever queries G or H with the decryption of c_{PKE}^* , B will be able to detect it using
 207 the plaintext-checking oracle. From where B is guaranteed to win the OW-PCA game.
 208 Therefore:

$$\begin{aligned} P[\text{QUERY } G^*] &\leq \text{Adv}_{\text{OW-PCA}}(B) \\ P[\text{QUERY } H^*] &\leq \text{Adv}_{\text{OW-PCA}}(B) \end{aligned}$$

209 Combining all inequalities above, we have:

$$\text{Adv}_0(A) \leq q_D \cdot \epsilon_{\text{MAC}} + 2\text{Adv}_{\text{OW-PCA}}(B)$$

210

□

211 4 Application to Kyber

212 CRYSTALS-Kyber [BDK⁺18][ABD⁺19] is an IND-CCA2 secure key encapsulation mech-
 213 anism that first constructs an IND-CPA secure public key encryption scheme whose
 214 security is based on the Module Learning with Error Problem (MLWE), then applies a
 215 generic transformation using *de-randomization* and *re-encryption* [HHK17]. The use of
 216 *re-encryption* for providing rigidity means that the decapsulation routine needs to run
 217 the encryption routine for verifying ciphertext integrity. Unfortunately for Kyber, the
 218 encryption routine consumes both more CPU cycles and carries a larger memory footprint.
 219 Therefore, applying an alternate transformation that replaces *re-encryption* with a MAC
 220 verification will bring substantial performance enhancement. This is particularly applicable
 221 to communication protocols such as TLS 1.3, where clients, often constrained environments,
 222 need to run the decapsulation routine.

223 The IND-CPA PKE of Kyber (algorithms 4, 5, 6 in [ABD⁺19]) is not OW-PCA secure. A
 224 plaintext-checking attack [RRCB19] can recover the secret key using a few maliciously
 225 constructed plaintext-checking queries for each coefficient of a Kyber secret key (there are
 226 512, 768, and 1024 coefficients in the secret key depending on the desired security level).
 227 However, we propose mitigation for plaintext-checking attack at protocol level by requiring
 228 each keypair to be used for decryption only once. If a decryption fails, then the secret
 229 key should be discarded and the key exchange terminated, and a new keypair should be
 230 generated to restart the key exchange. In fact, such ephemeral key exchange is already
 231 required in TLS 1.3 for forward secrecy, so requiring the one-time use of keypair does not
 232 introduce additional operational cost for protocols such as TLS 1.3.

233 The routines of CCAKEM can be found in Algorithm 7, 8, 9 in [ABD⁺19]. We modify
 234 Algorithms 8 and 9 using authenticated encryption (AE) mode as follows where Algorithm
 235 7 is unchanged.

Algorithm 12 Kyber.CCAKEM.KeyGen()

- 1: $z \xleftarrow{\$} \mathcal{B}^{32}$
 - 2: $(\text{pk}, \text{sk}') \xleftarrow{\$} \text{Kyber.CPAPKE.KeyGen}()$
 - 3: $\text{sk} = (\text{sk}', \text{pk}, H(\text{pk}), z)$ ▷ H is instantiated with SHA3-256
 - 4: **return** (pk, sk)
-

Algorithm 13 $\text{Kyber.CCAKEM.Encap}^+(\text{pk})$

```

1:  $m \xleftarrow{\$} \mathcal{B}^{32}$ 
2:  $m' = H(m)$ 
3:  $(\bar{K}, K_{\text{MAC}}) = G(m' \| H(\text{pk}))$ 
4:  $r \xleftarrow{\$} \mathcal{R}$ 
5:  $c' \leftarrow \text{Kyber.CPAPKE.Enc}(\text{pk}, m', r)$ 
6:  $t = \text{MAC}(K_{\text{MAC}}, c')$ 
7:  $K = \text{KDF}(\bar{K} \| t)$ 
8:  $c \leftarrow (c', t)$ 
9: return  $(c, K)$ 

```

\triangleright Do not output system RNG directly
 \triangleright G is instantiated with SHA3-512
 \triangleright KDF is instantiated with Shake256

Algorithm 14 $\text{Kyber.CCAKEM.Decap}^+(\text{sk}, c)$ **Require:** Secret key $\text{sk} = (\text{sk}', \text{pk}, H(\text{pk}), z)$ **Require:** Ciphertext $c = (c', t)$

```

1:  $(\text{sk}', \text{pk}, h, z) \leftarrow \text{sk}$ 
2:  $(c', t_1) \leftarrow c$ 
3:  $\hat{m} = \text{Kyber.CPAPKE.Dec}(\text{sk}', c')$ 
4:  $(\bar{K}, K_{\text{MAC}}) = G(\hat{m}' \| h)$ 
5:  $\hat{t} = \text{MAC}(K_{\text{MAC}}, c)$ 
6: if  $\hat{t} = t$  then
7:    $K = \text{KDF}(\bar{K} \| t)$ 
8: else
9:    $K = \text{KDF}(z \| t)$ 
10: end if
11: return  $K$ 

```

236 *Remark 1.* If c is manipulated, then the verification of t_1 will be failed. In this case, there
237 is no K outputted from the decap^+ . So the attacks described in the following subsections
238 won't work. We also added the key confirmation, which is tag t_2 .

239 Note for authenticated encryption, tags t_i 's are necessary for the inputs.

240 In fact, this is authenticated encryption instead of EtM. So we should change that,
241 called authenticated encryption. Please do not change my notation. They have their
242 meanings in AE mode.

243 4.1 MAC performance

244 When instantiating an instance $\text{KEM}(\text{KeyGen}, \text{Encap}, \text{Decap})$, there are a variety of possible
245 MAC's to choose from. For each of the security level and choice of MAC, we measured the
246 number of CPU cycles needed to produce a digest of the unauthenticated ciphertext. The
247 median (top) and mean (bottom) measurements are reported in table 1.

Table 1: Standalone MAC performances

Name	Security	measurement	768 bytes	1088 bytes	1568 bytes
CMAC	many-time	median	5022	5442	6090
		mean	5131	5578	6154
KMAC-256	many-time	median	7934	9862	11742
		mean	8594	10693	12319
GMAC	one-time	median	2778	2756	2762
		mean	2843	2780	2919
Poly1305	one-time	median	1128	1218	1338
		mean	1435	1504	1625

Based on the security reduction in section 3 we chose standalone GMAC and Poly1305 for their substantial performance advantage over other constructions such as CBC-MAC and KMAC. We then modified the reference implementation (TODO: citation needed <https://github.com/pq-crystals/kyber>) according to algorithms 12, 13, and 14. We measured the median and mean CPU cycles needed to run each of the routines. The measurements are listed in table 2

Table 2: Kyber-AE performance measurements

Name	Security level	measurement	KeyGen	Encap	Decap
Kyber512	128 bits	median			
		mean			
Kyber768	192 bits	median			
		mean			
Kyber1024	256 bits	median			
		mean			
KyberAE512 w/ GMAC	128 bits	median			
		mean			
KyberAE768 w/ GMAC	192 bits	median			
		mean			
KyberAE1024 w/ GMAC	256 bits	median			
		mean			
KyberAE512 w/ Poly1305	128 bits	median			
		mean			
KyberAE768 w/ Poly1305	192 bits	median			
		mean			
KyberAE1024 w/ Poly1305	256 bits	median			
		mean			

5 Conclusions and future works

References

- [ABD⁺19] Roberto Avanzi, Joppe Bos, Léo Ducas, Eike Kiltz, Tancrede Lepoint, Vadim Lyubashevsky, John M Schanck, Peter Schwabe, Gregor Seiler, and Damien Stehlé. Crystals-kyber algorithm specifications and supporting documentation. *NIST PQC Round, 2*(4):1–43, 2019.
- [BCD⁺16] Joppe Bos, Craig Costello, Léo Ducas, Ilya Mironov, Michael Naehrig, Valeria Nikolaenko, Ananth Raghunathan, and Douglas Stebila. Frodo: Take off the

- ring! practical, quantum-secure key exchange from lwe. In *Proceedings of the 2016 ACM SIGSAC conference on computer and communications security*, pages 1006–1018, 2016.
- [BDK⁺18] Joppe Bos, Léo Ducas, Eike Kiltz, Tancrede Lepoint, Vadim Lyubashevsky, John M Schanck, Peter Schwabe, Gregor Seiler, and Damien Stehlé. Crystals-kyber: a cca-secure module-lattice-based kem. In *2018 IEEE European Symposium on Security and Privacy (EuroS&P)*, pages 353–367. IEEE, 2018.
- [BN00] Mihir Bellare and Chanathip Namprempre. Authenticated encryption: Relations among notions and analysis of the generic composition paradigm. In *International Conference on the Theory and Application of Cryptology and Information Security*, pages 531–545. Springer, 2000.
- [BP18] Daniel J Bernstein and Edoardo Persichetti. Towards kem unification. *Cryptology ePrint Archive*, 2018.
- [BS20] Dan Boneh and Victor Shoup. A graduate course in applied cryptography. *Draft 0.5*, 2020.
- [DKSRV18] Jan-Pieter D’Anvers, Angshuman Karmakar, Sujoy Sinha Roy, and Frederik Vercauteren. Saber: Module-lwr based key exchange, cpa-secure encryption and cca-secure kem. In *Progress in Cryptology–AFRICACRYPT 2018: 10th International Conference on Cryptology in Africa, Marrakesh, Morocco, May 7–9, 2018, Proceedings 10*, pages 282–305. Springer, 2018.
- [FO99] Eiichiro Fujisaki and Tatsuaki Okamoto. Secure integration of asymmetric and symmetric encryption schemes. In *Annual international cryptology conference*, pages 537–554. Springer, 1999.
- [HHK17] Dennis Hofheinz, Kathrin Hövelmanns, and Eike Kiltz. A modular analysis of the fujisaki-okamoto transformation. In *Theory of Cryptography Conference*, pages 341–371. Springer, 2017.
- [KE23] NIST Module-Lattice-Based Key-Encapsulation. Mechanism standard. *NIST Post-Quantum Cryptography Standardization Process; NIST: Gaithersburg, MD, USA*, 2023.
- [RRCB19] Prasanna Ravi, Sujoy Sinha Roy, Anupam Chattopadhyay, and Shivam Bhasin. Generic side-channel attacks on cca-secure lattice-based pke and kem schemes. *Cryptology ePrint Archive*, 2019.
- [Sho04] Victor Shoup. Sequences of games: a tool for taming complexity in security proofs. *cryptology eprint archive*, 2004.
- [UXT⁺22] Rei Ueno, Keita Xagawa, Yutaro Tanaka, Akira Ito, Junko Takahashi, and Naofumi Homma. Curse of re-encryption: a generic power/em analysis on post-quantum kems. *IACR Transactions on Cryptographic Hardware and Embedded Systems*, pages 296–322, 2022.